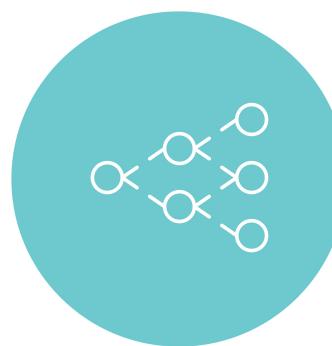


Docker Training



Pragra
ONWARD & UPWARD

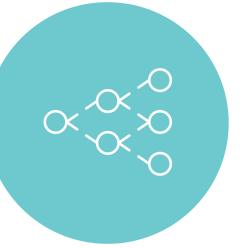


Welcome

[www.docker.com]

Agenda

- Understanding Docker
- Docker Containers
- Docker Images
- Docker Networking
- Docker volume
- Docker Compose

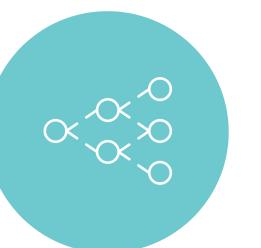


What is Docker

Docker is an open-source project that automates the deployment of applications inside software containers, by providing an additional layer of abstraction and automation of operating system-level virtualization on Linux

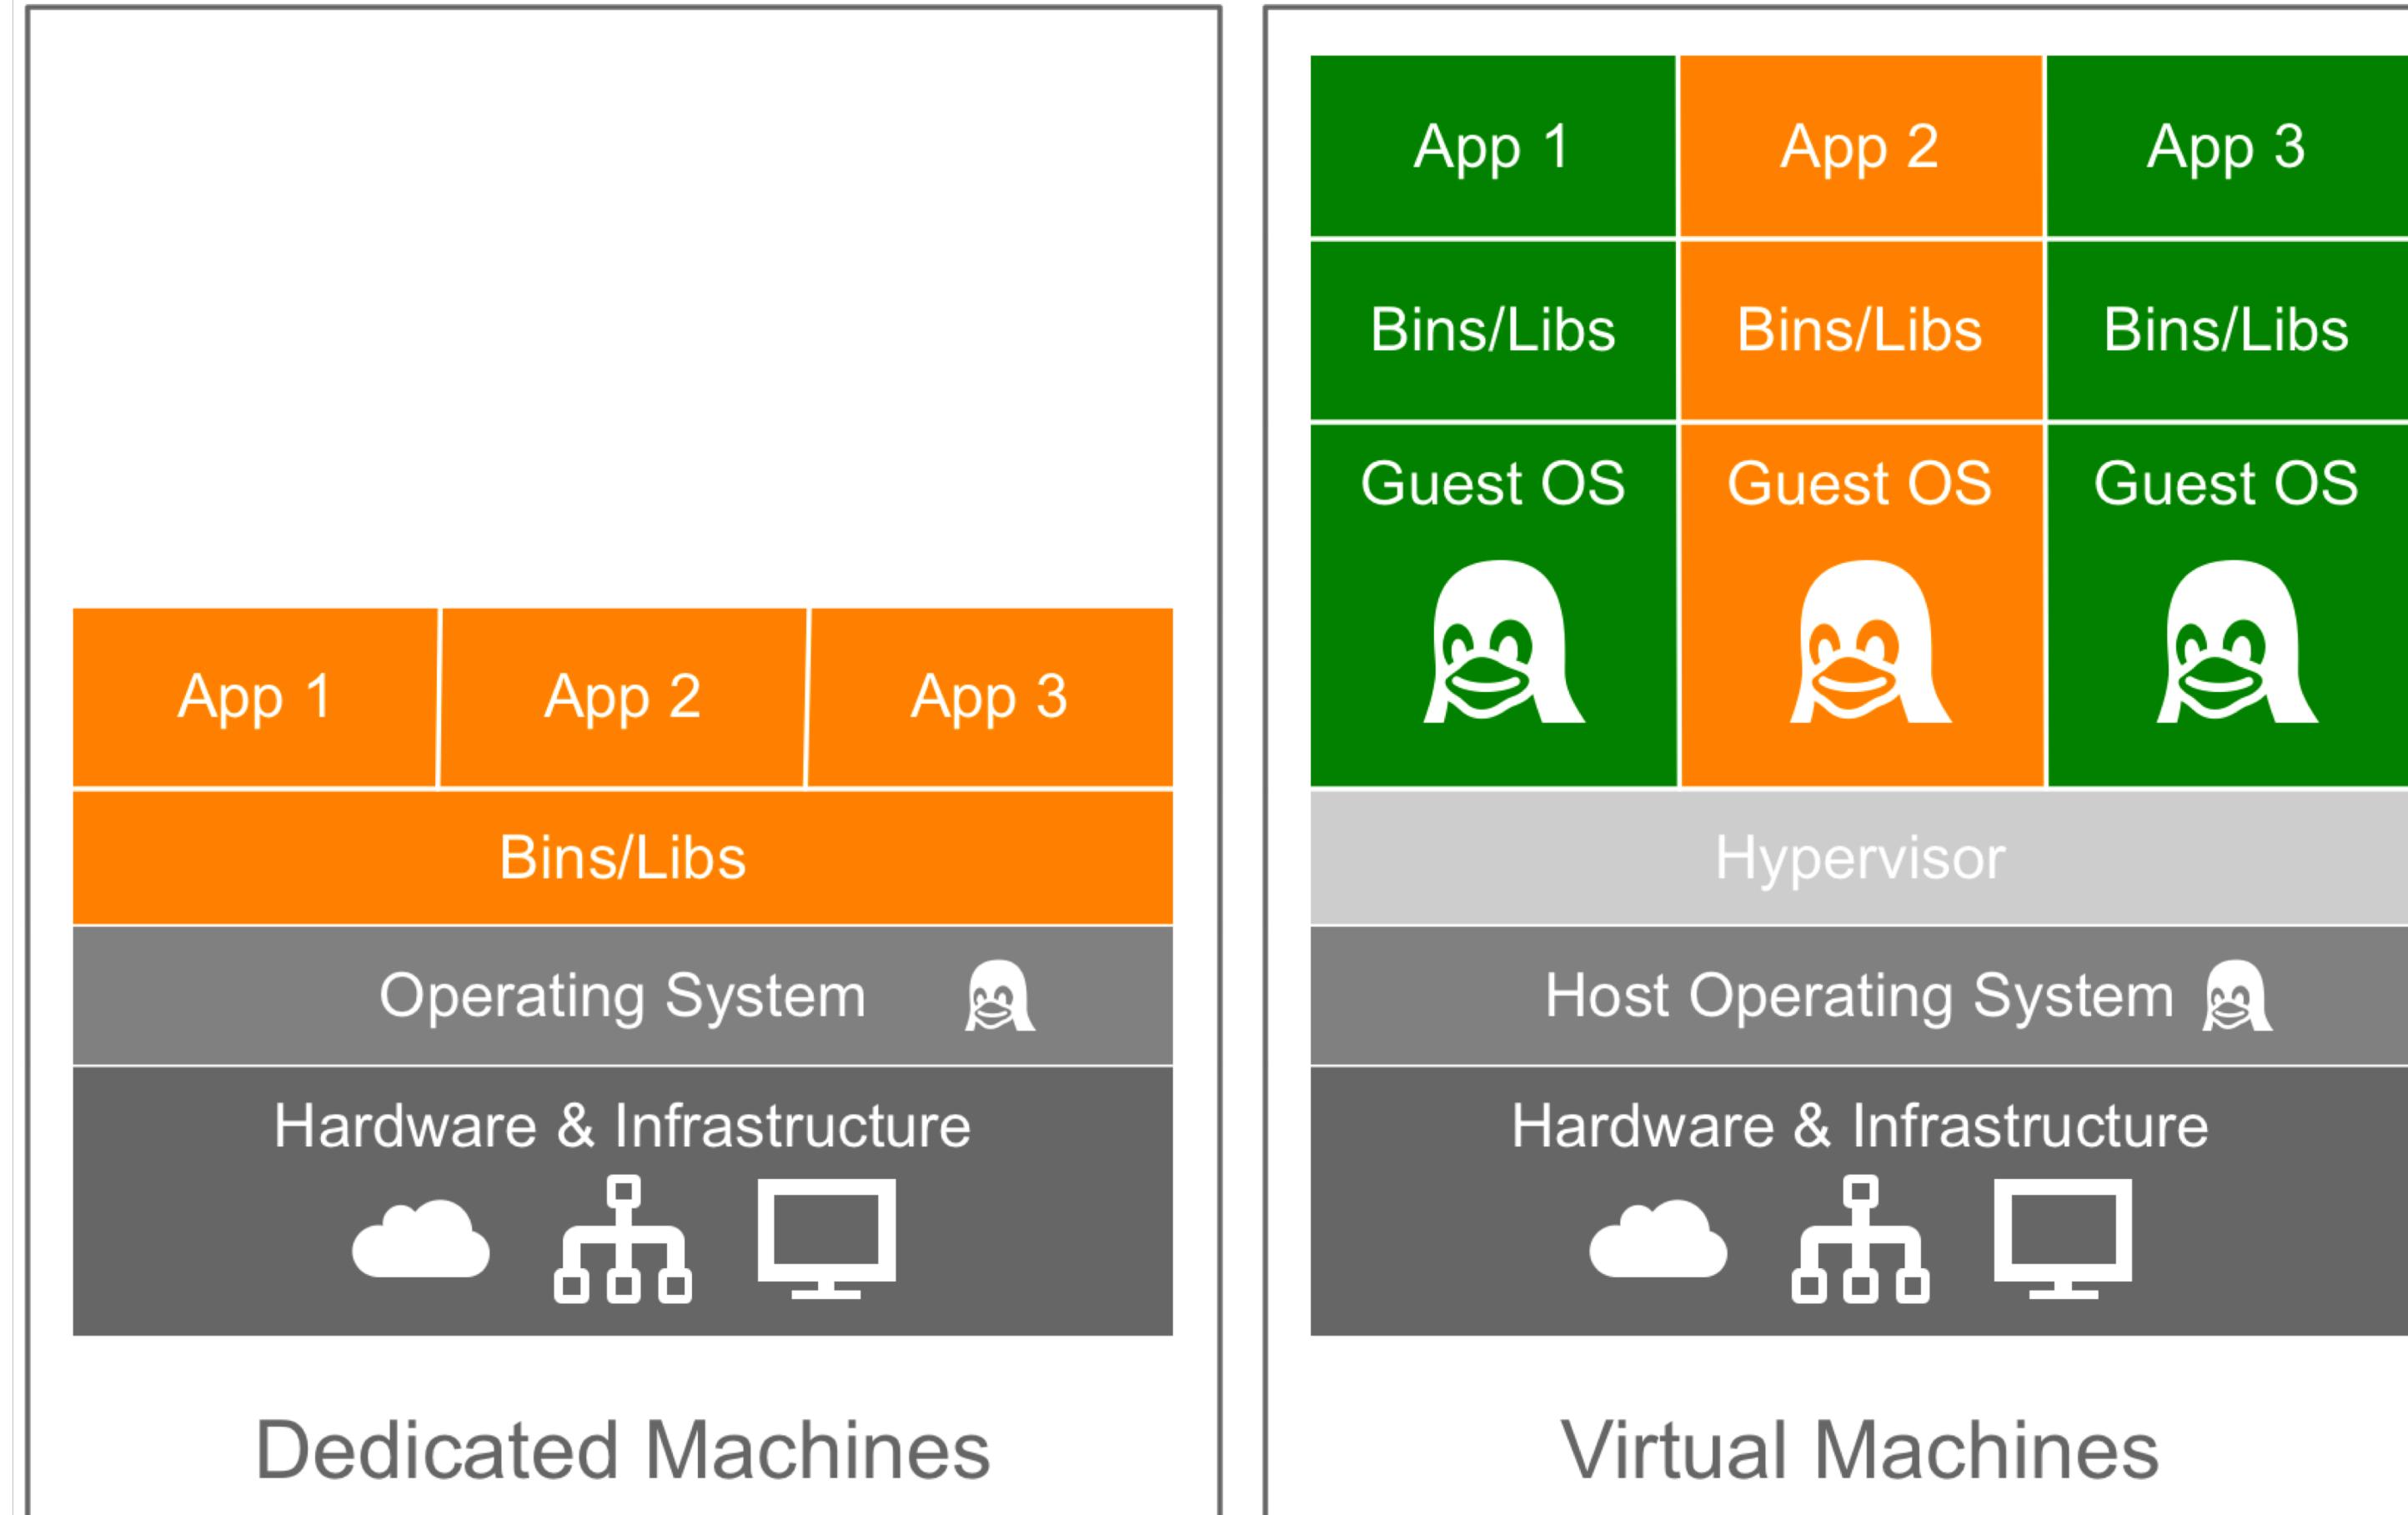
Provide a uniformed wrapper around a software package: «*Build, Ship and Run Any App, Anywhere*» [\[www.docker.com\]](http://www.docker.com)

Similar to shipping containers: The container is always the same, regardless of the contents and thus fits on all trucks, cranes, ships, ...

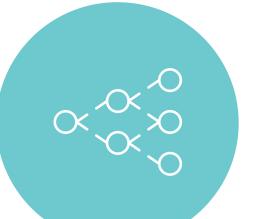


Difference Between Host/VM and Docker

- Docker is a container management system that helps us easily manage **Linux Containers (LXC)** in an easier and universal fashion.



Unlike VM docker doesn't install guest Operating system with own kernel. It used host Kernel and provide Isolation at binary level



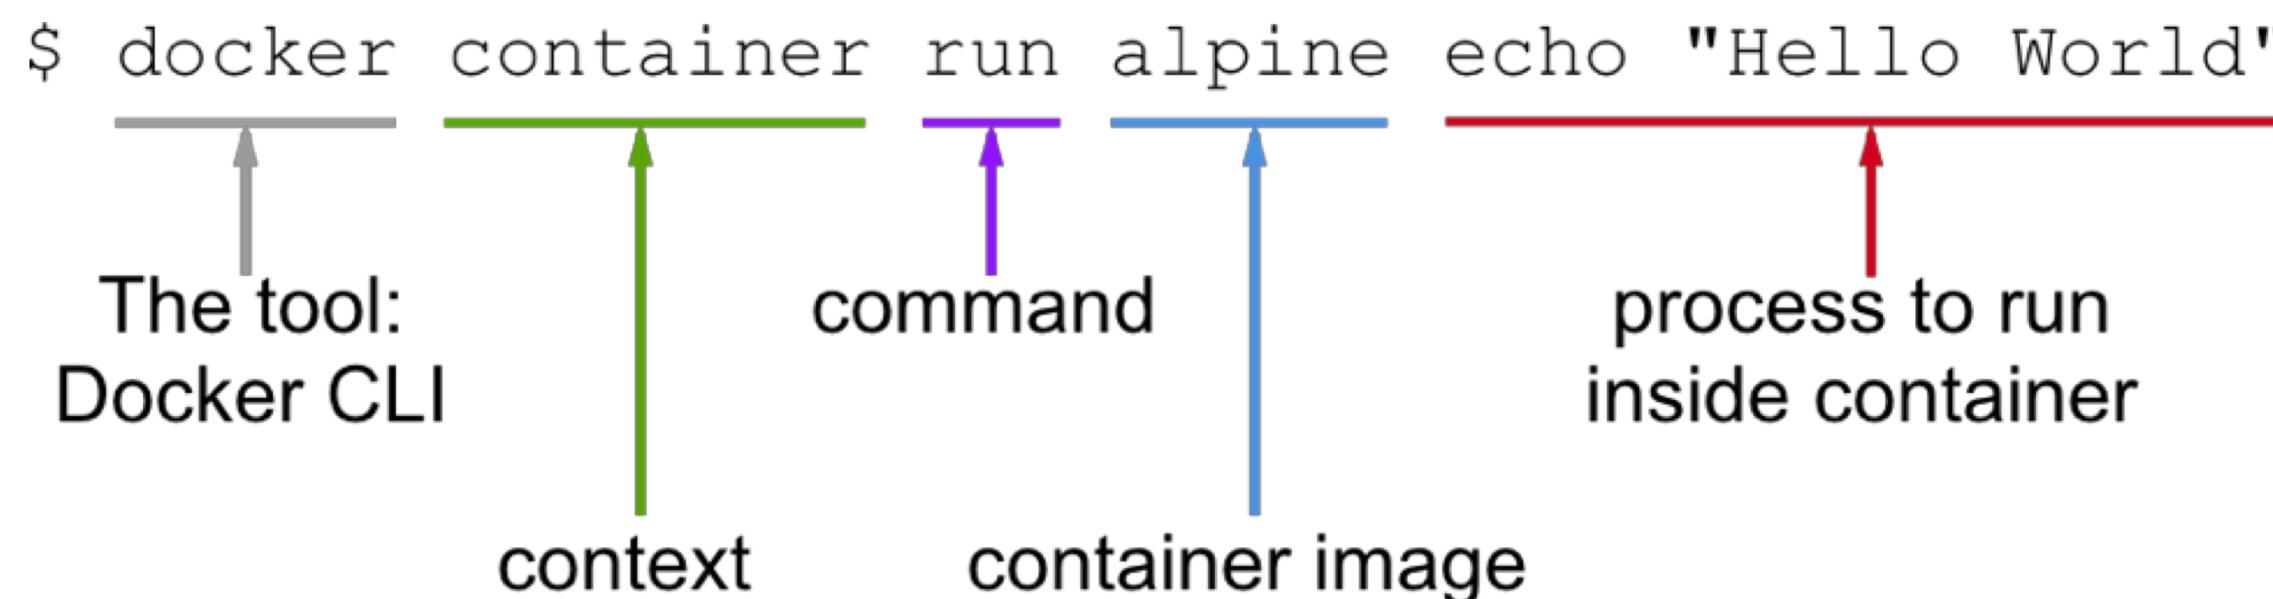
Docker CLI

- Docker comes with sophisticated CLI - docker cli provide various management command.

Check version of your docker installation

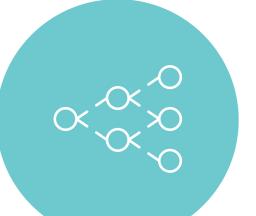
→ ~ docker -v

Docker version 18.09.2, build 6247962



Management Commands:

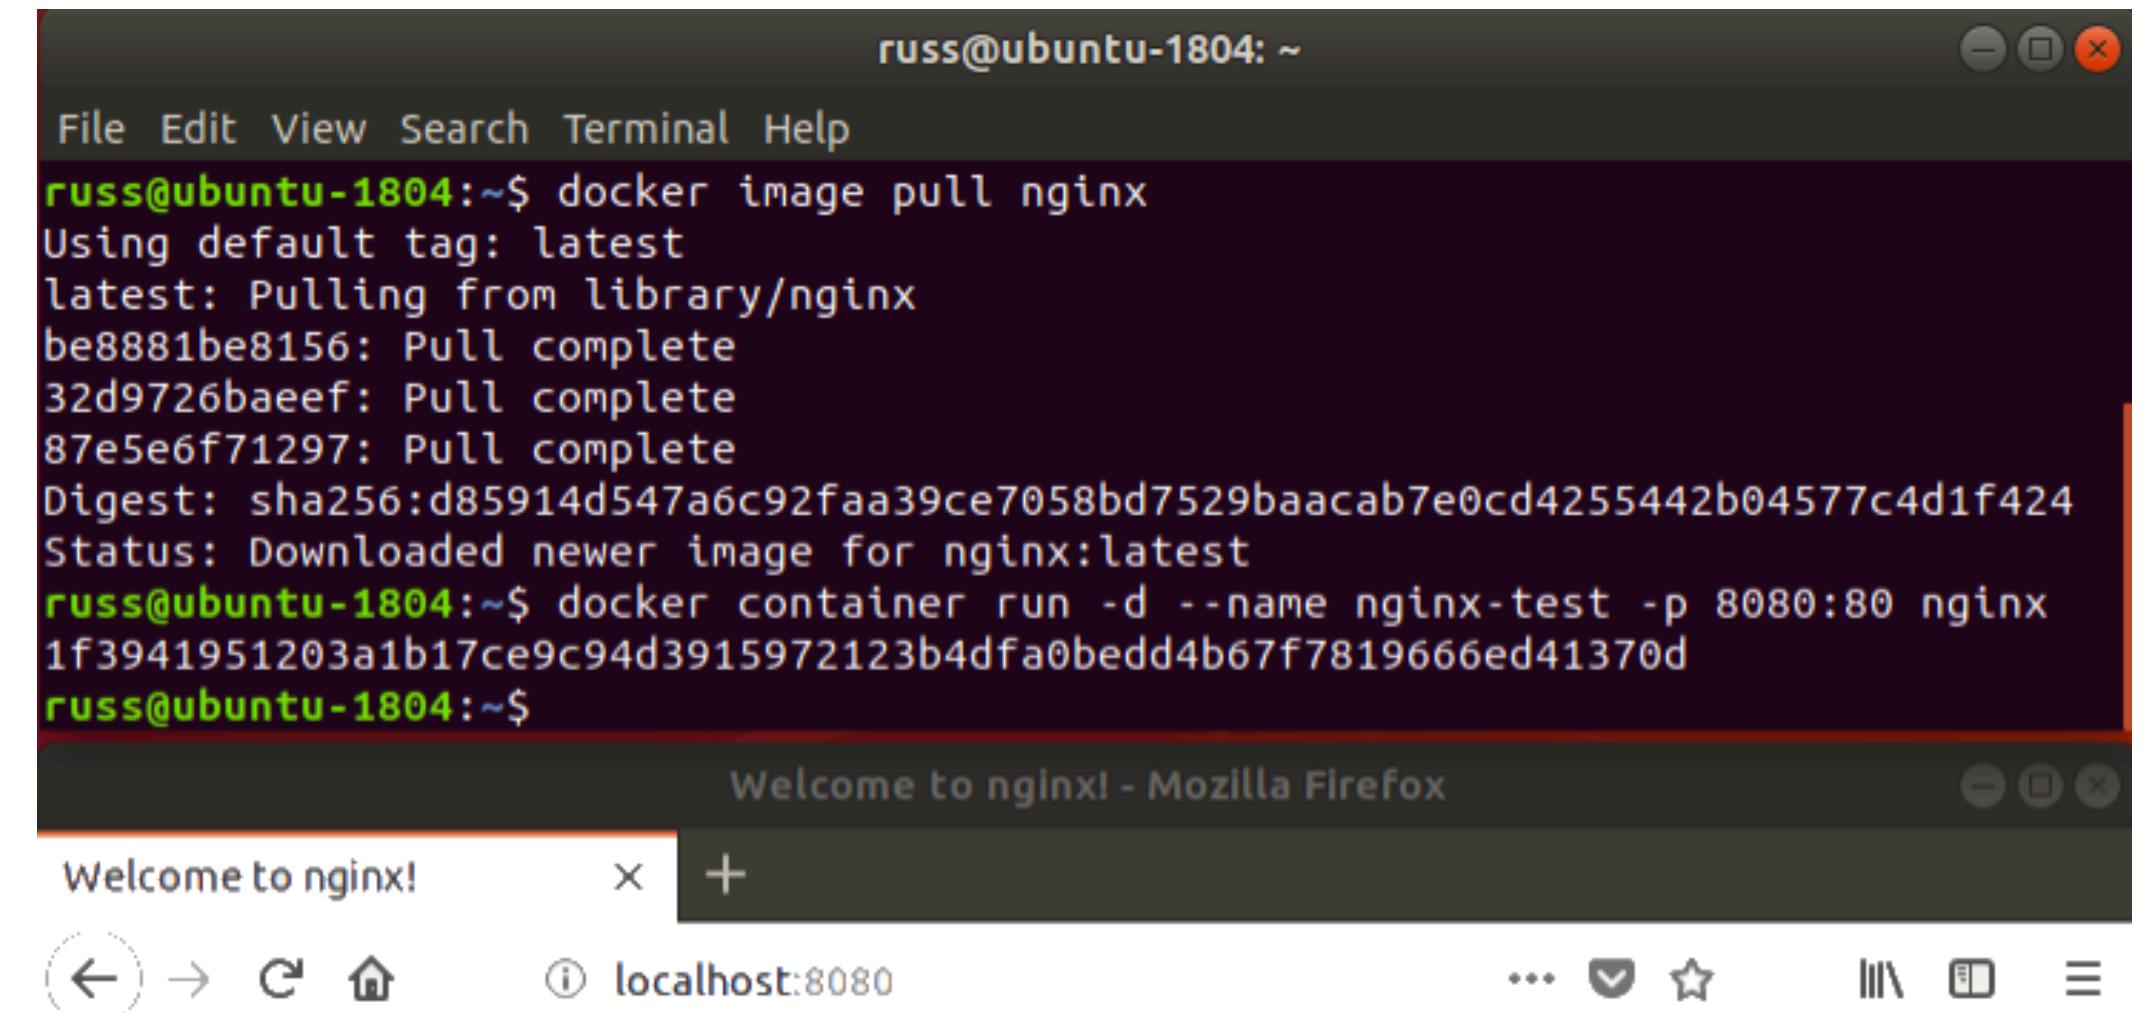
builder	Manage builds
config	Manage Docker configs
container	Manage containers
image	Manage images
network	Manage networks
node	Manage Swarm nodes
plugin	Manage plugins
secret	Manage Docker secrets
service	Manage services
stack	Manage Docker stacks
swarm	Manage Swarm
system	Manage Docker
volume	Manage volumes



Run your first Container

```
~ docker container run -d --name nginx-test -p 8080:80 nginx
```

- Command will run the a container name nginx-test and will forward port 80 from container to 8080 of localhost



The terminal window shows the following commands:

```
russ@ubuntu-1804:~$ docker image pull nginx
Using default tag: latest
latest: Pulling from library/nginx
be8881be8156: Pull complete
32d9726baeef: Pull complete
87e5e6f71297: Pull complete
Digest: sha256:d85914d547a6c92faa39ce7058bd7529baacab7e0cd4255442b04577c4d1f424
Status: Downloaded newer image for nginx:latest
russ@ubuntu-1804:~$ docker container run -d --name nginx-test -p 8080:80 nginx
1f3941951203a1b17ce9c94d3915972123b4dfa0bedd4b67f7819666ed41370d
russ@ubuntu-1804:~$
```

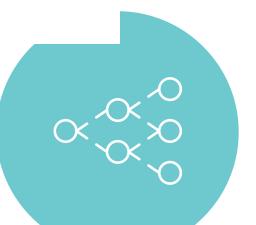
The browser window shows the nginx welcome page at localhost:8080.

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working.
Further configuration is required.

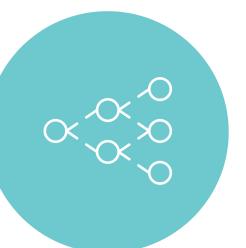
For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.



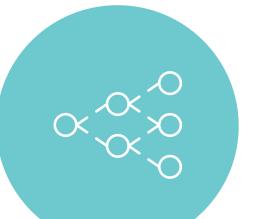
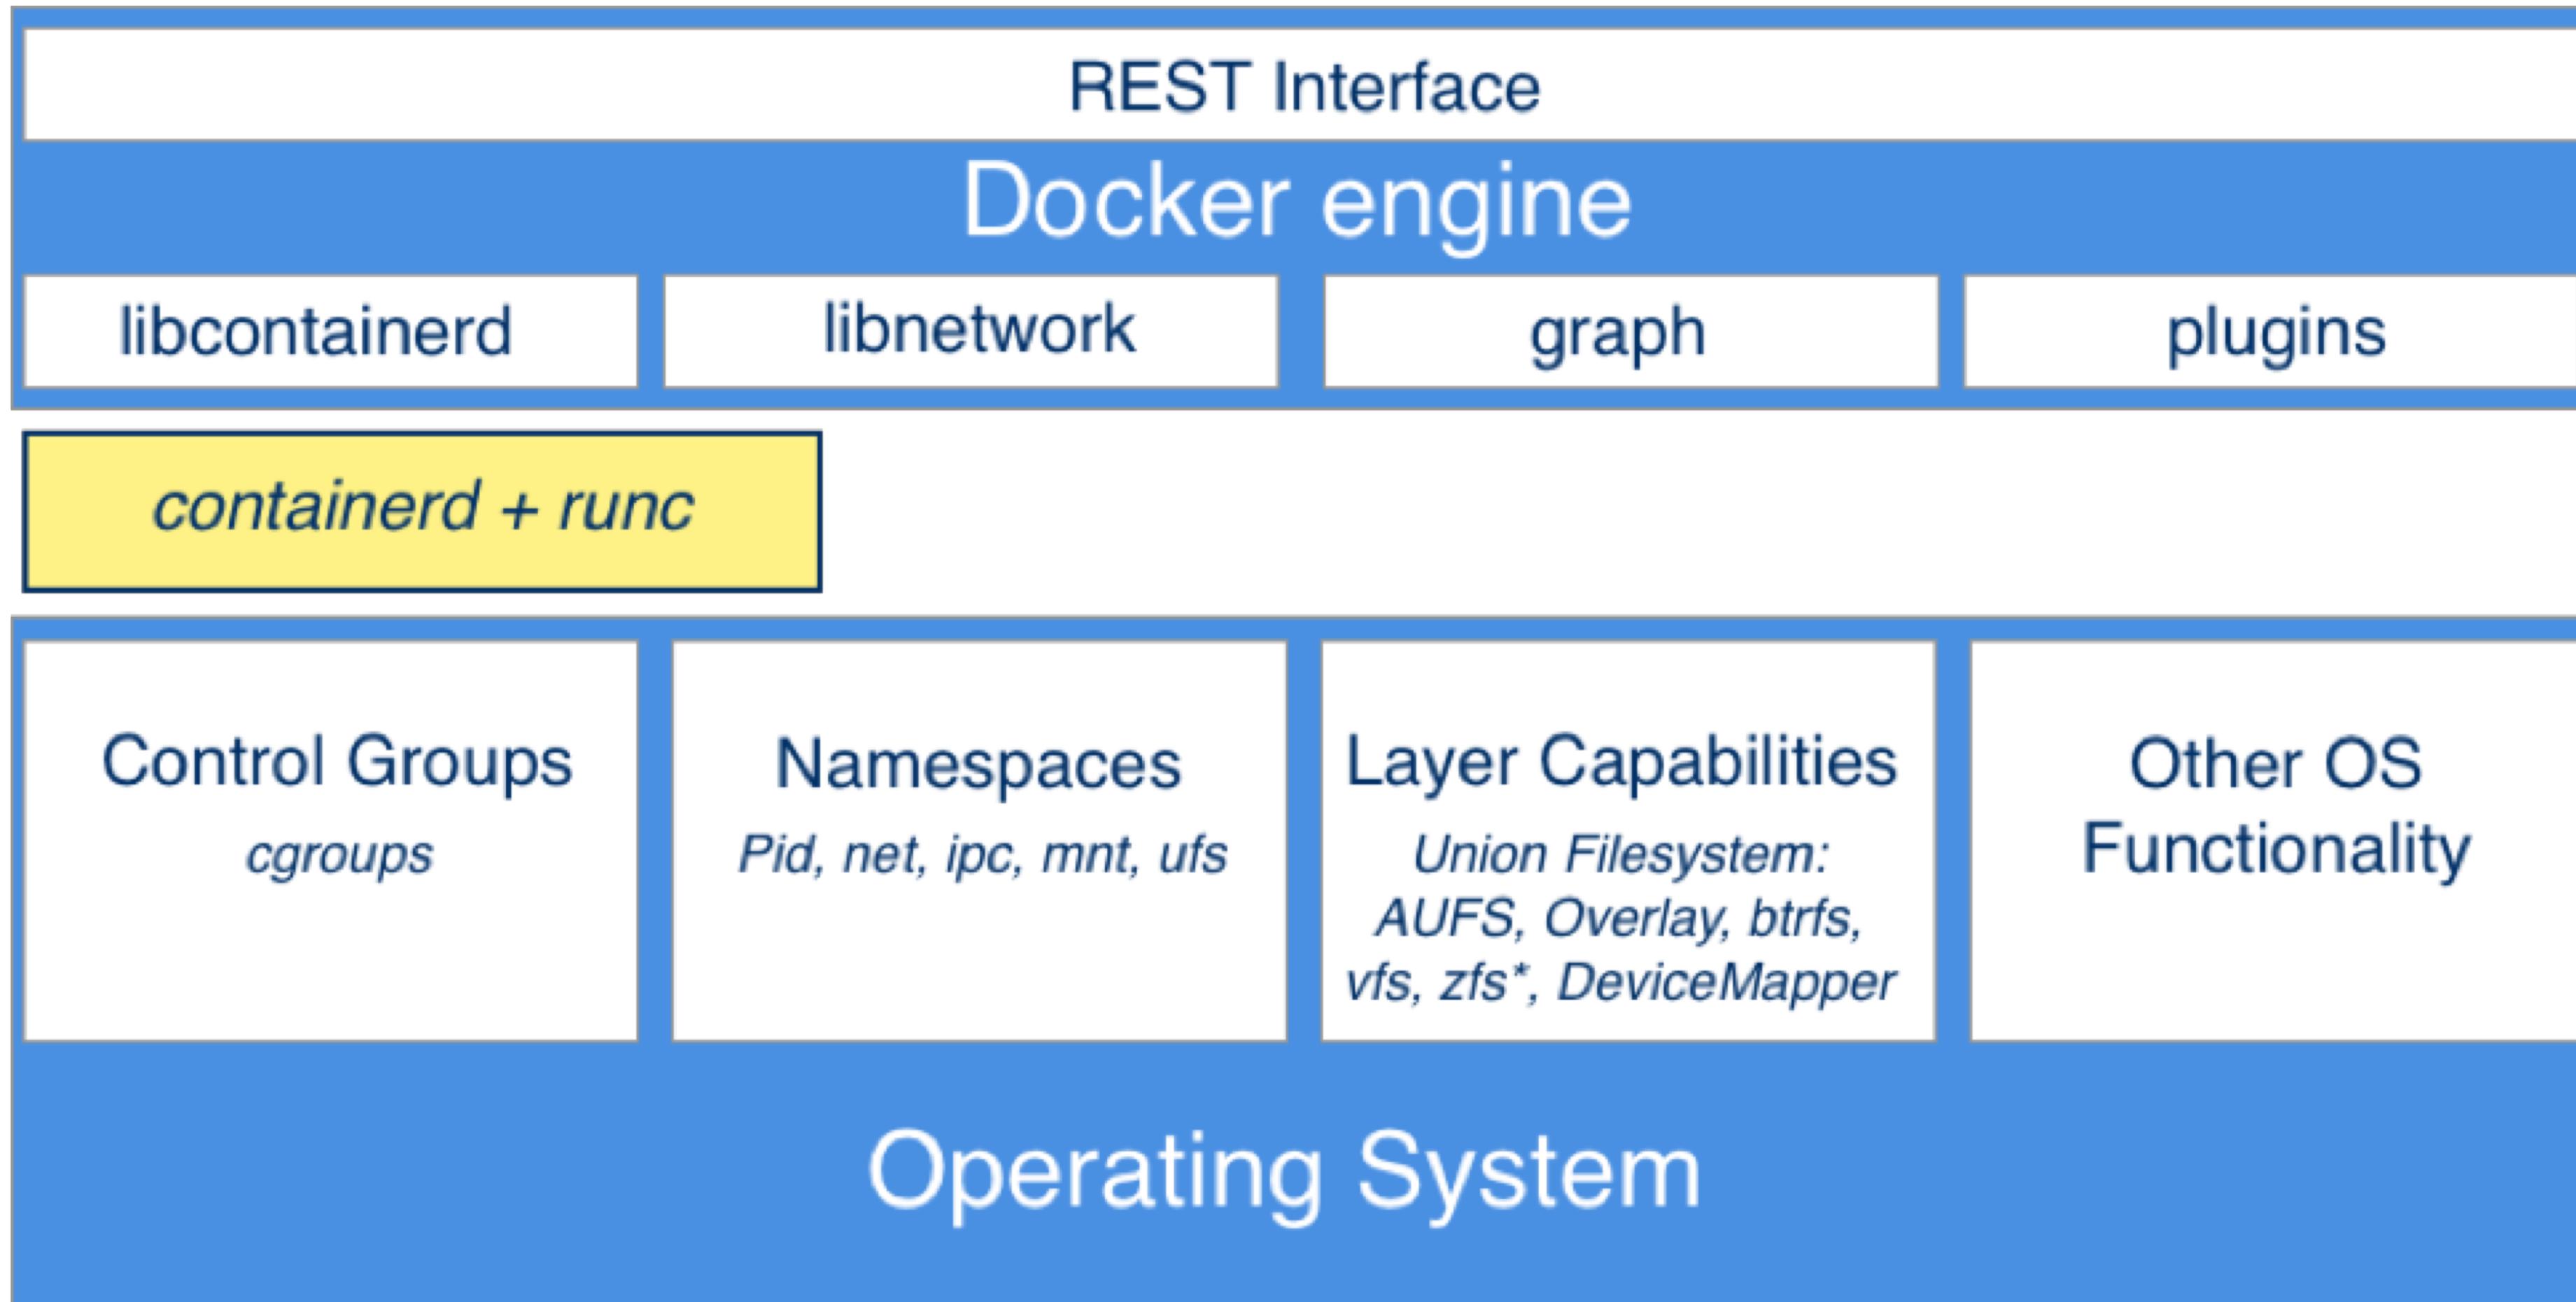
Questions ?

- What command did we use to run a container?
 - Can you run a docker in a VM
 - Can we have docker and VM on same machine ?
 - Why Docker is better than VM ?
-
- Docker CLI restructure blog post: <https://blog.docker.com/2017/01/whats-new-in-docker-1-13/>
 - Docker Extended Support Announcement: <https://blog.docker.com/2018/07/extending-support-cycle-docker-community-edition/>



Docker Architecture

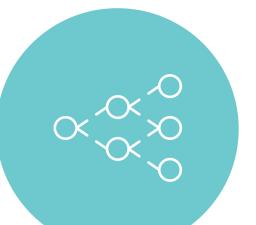
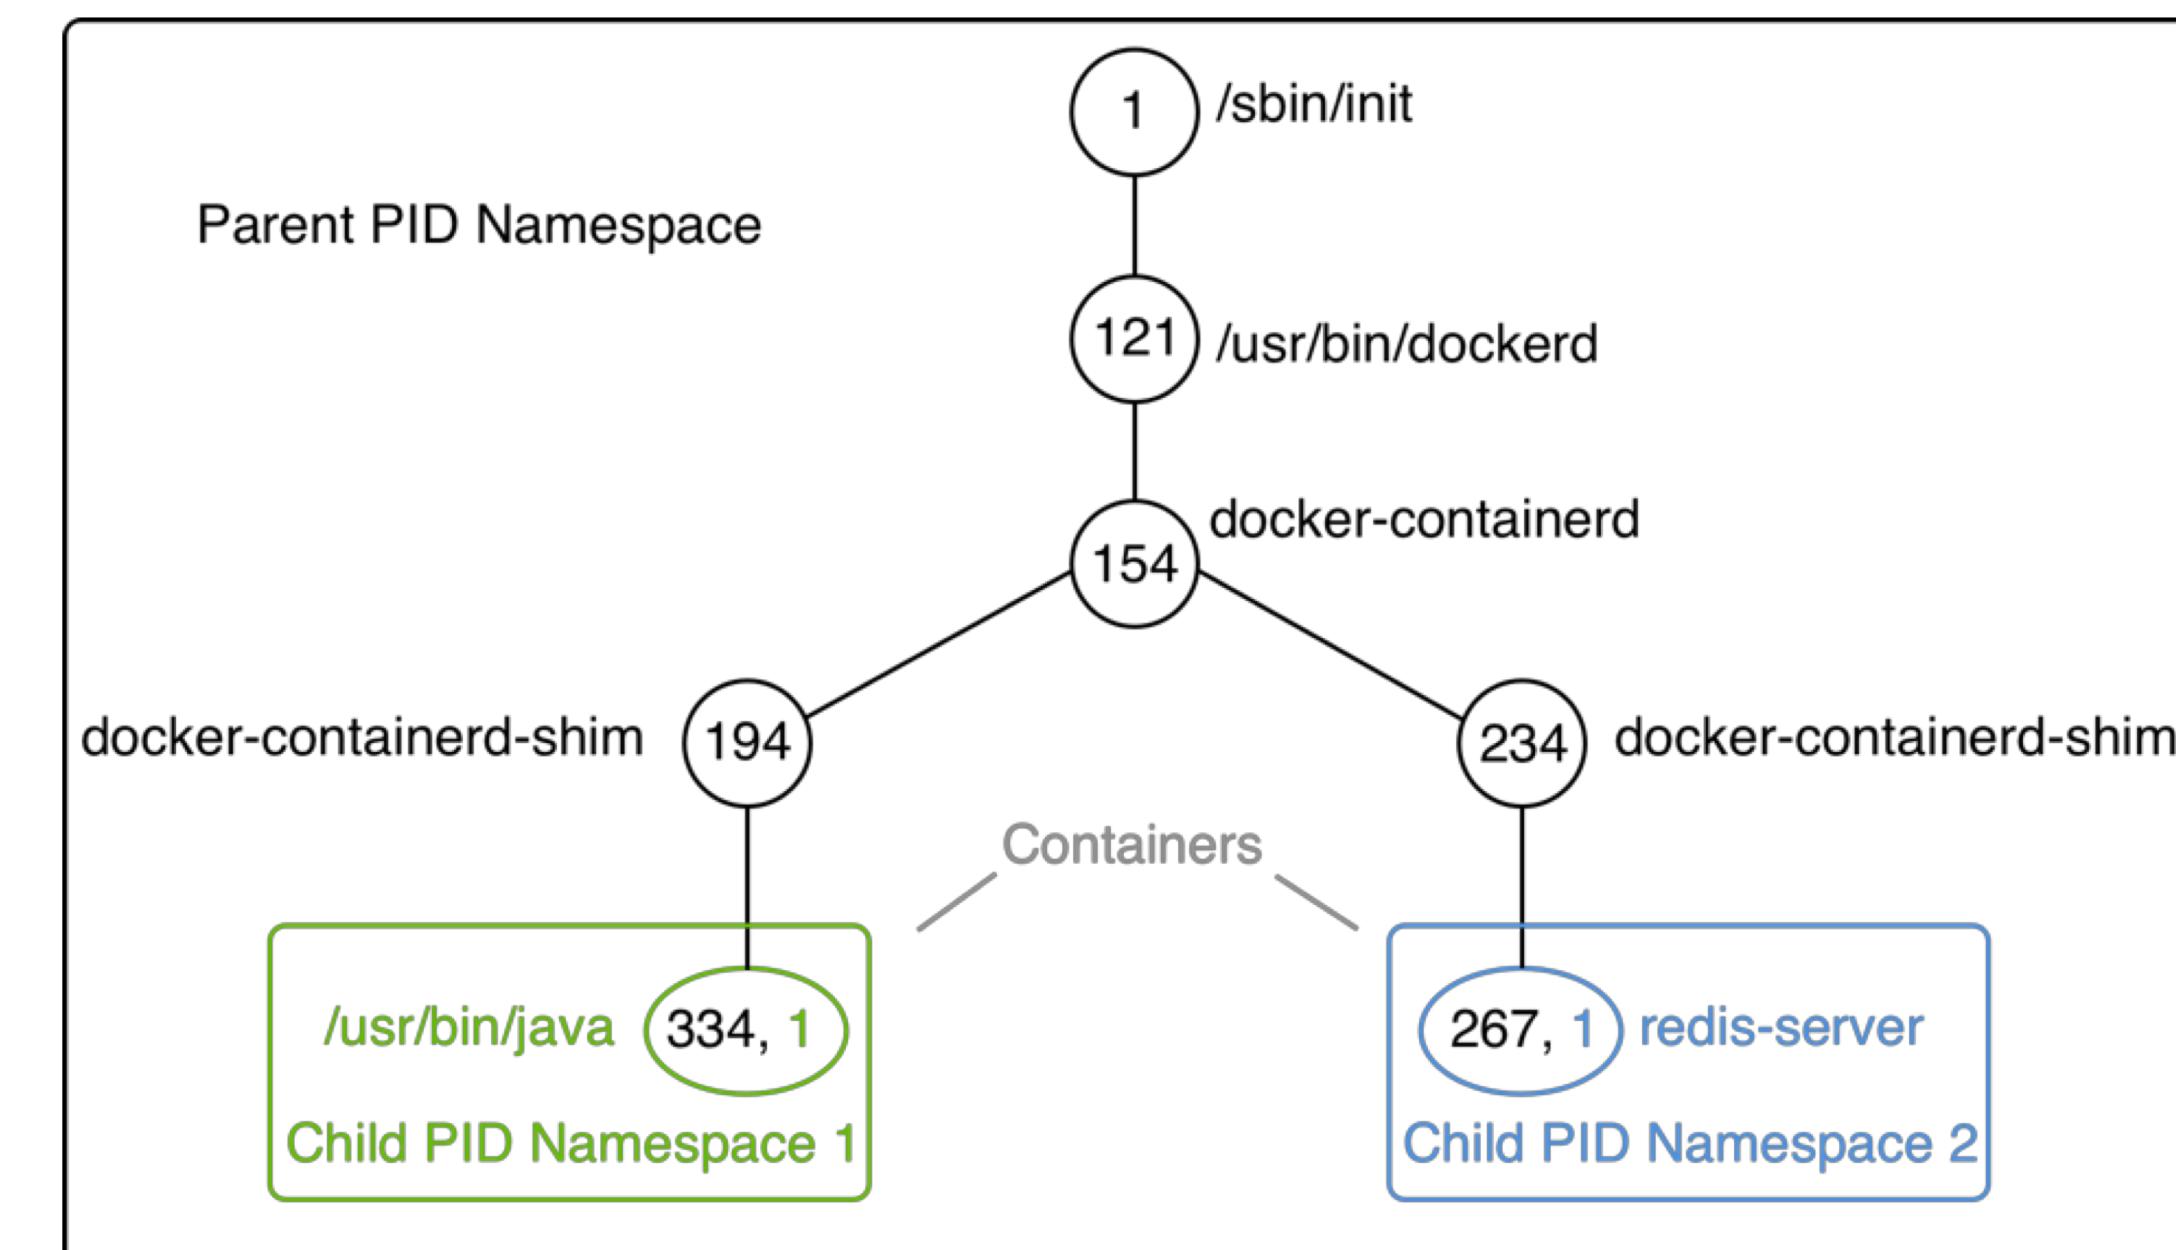
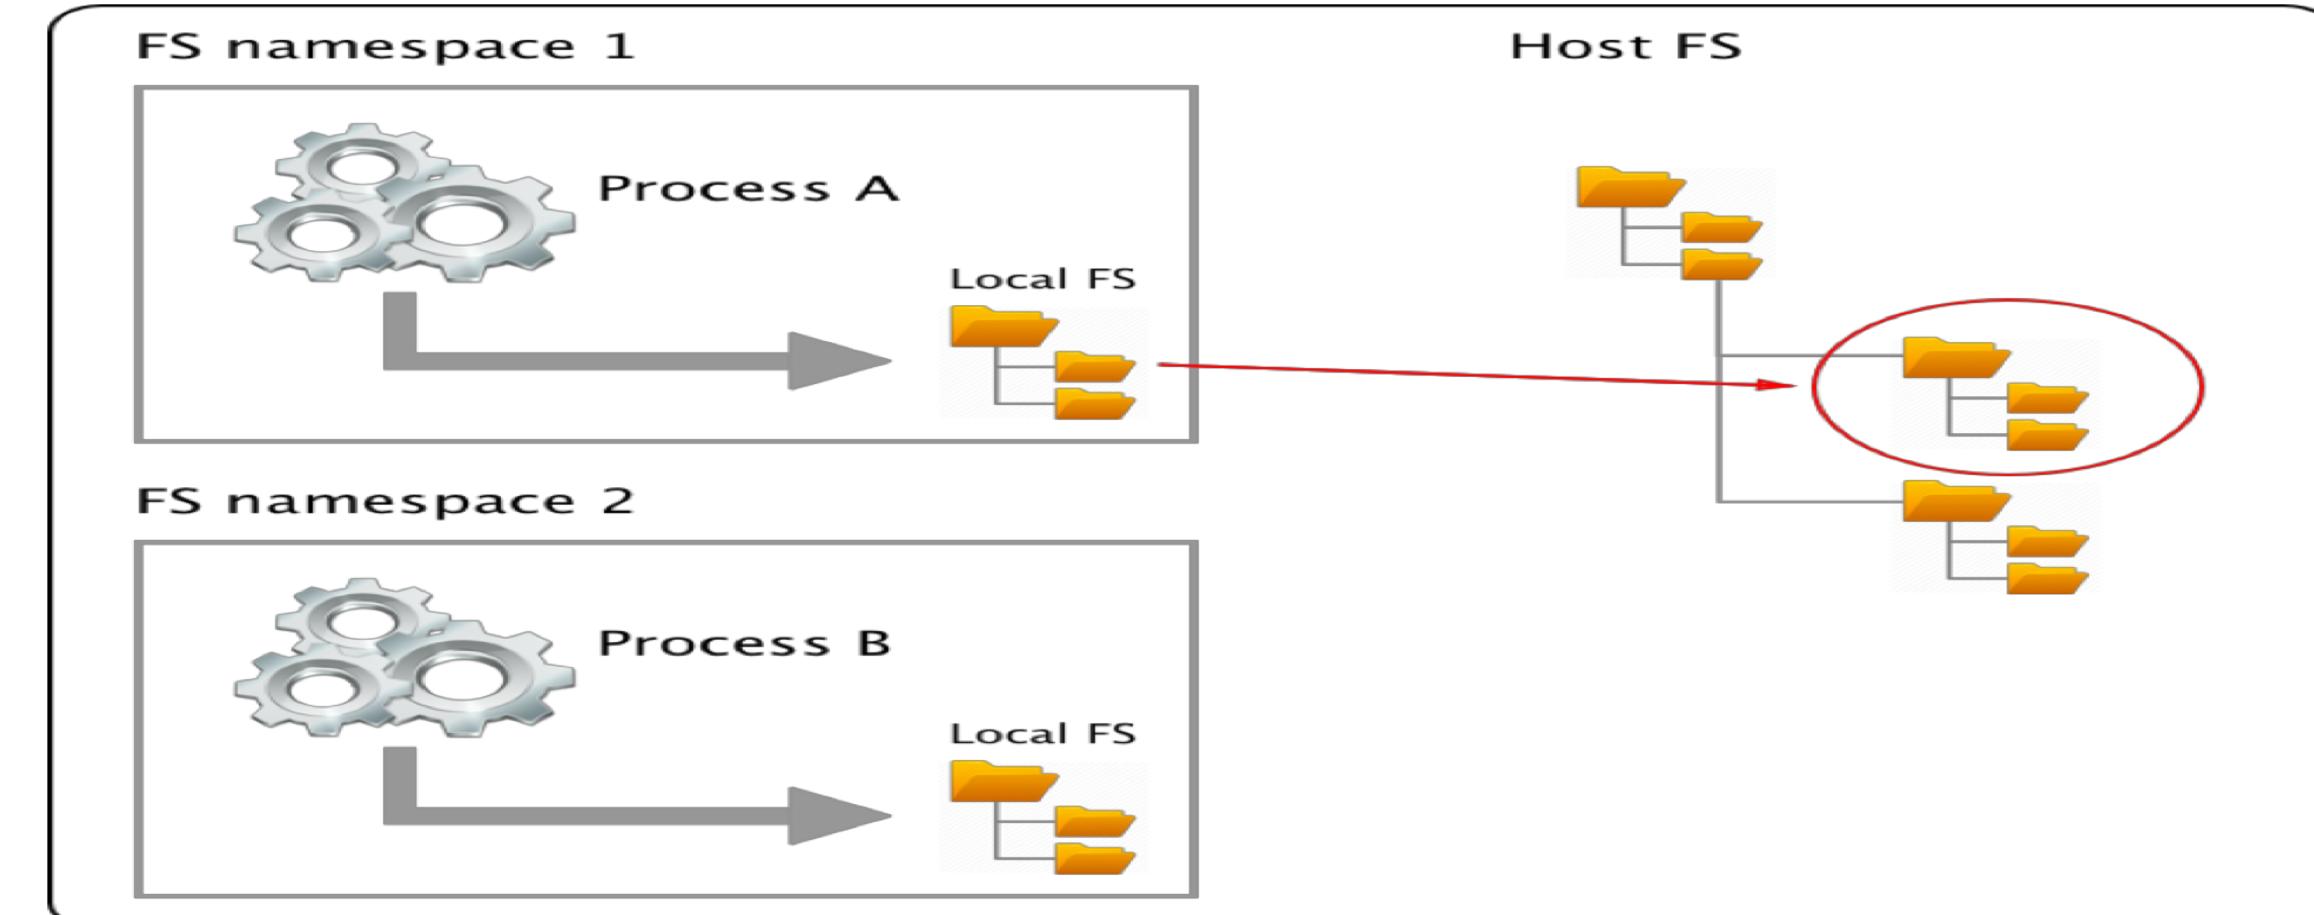
8



Linux Namespaces

A namespace is an abstraction of global resources such as filesystems, network access, process tree (also named PID namespace) or the system group IDs, and user IDs. A Linux system is initialized with a single instance of each namespace type. After initialization, additional namespaces can be created or joined.

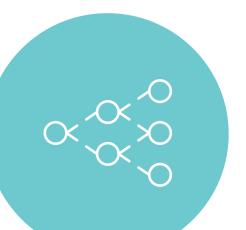
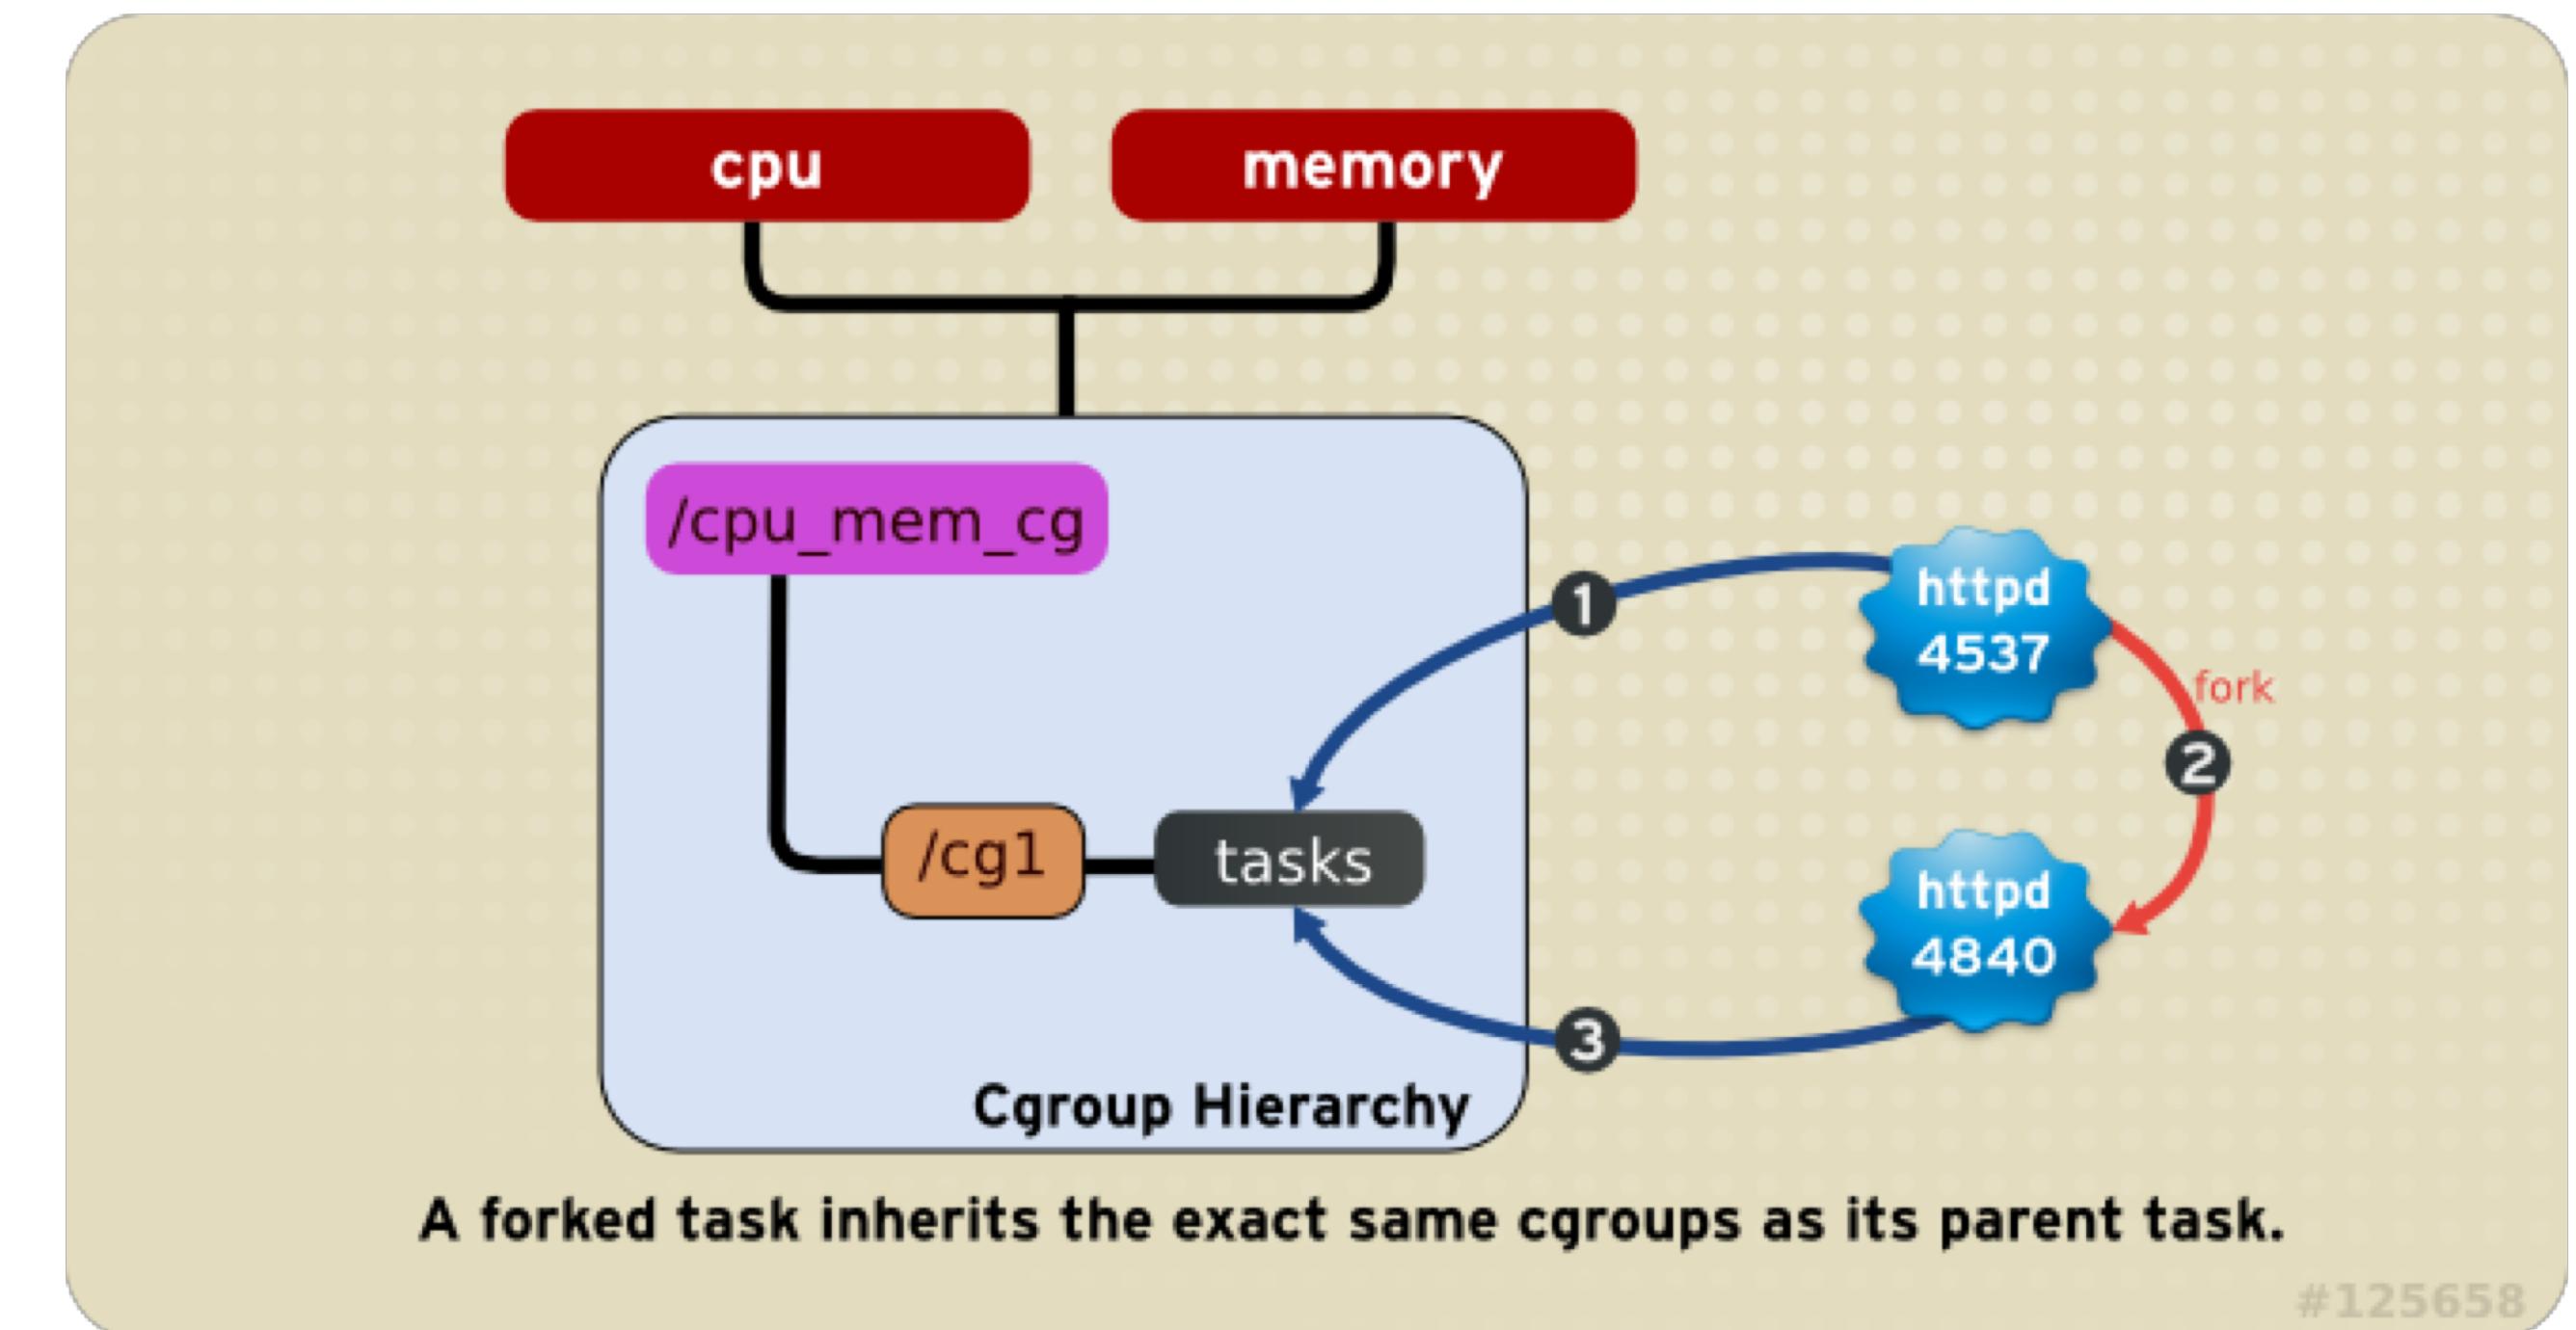
Docker creates child namespace for each container.



Control Groups (CGROUP)

10

Linux cgroups are used to limit, manage, and isolate resource usage of collections of processes running on a system. Resources are CPU time, system memory, network bandwidth, or combinations of these resources, and so on.

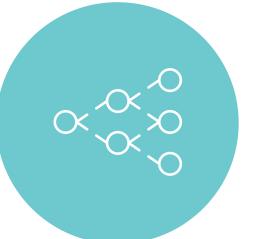
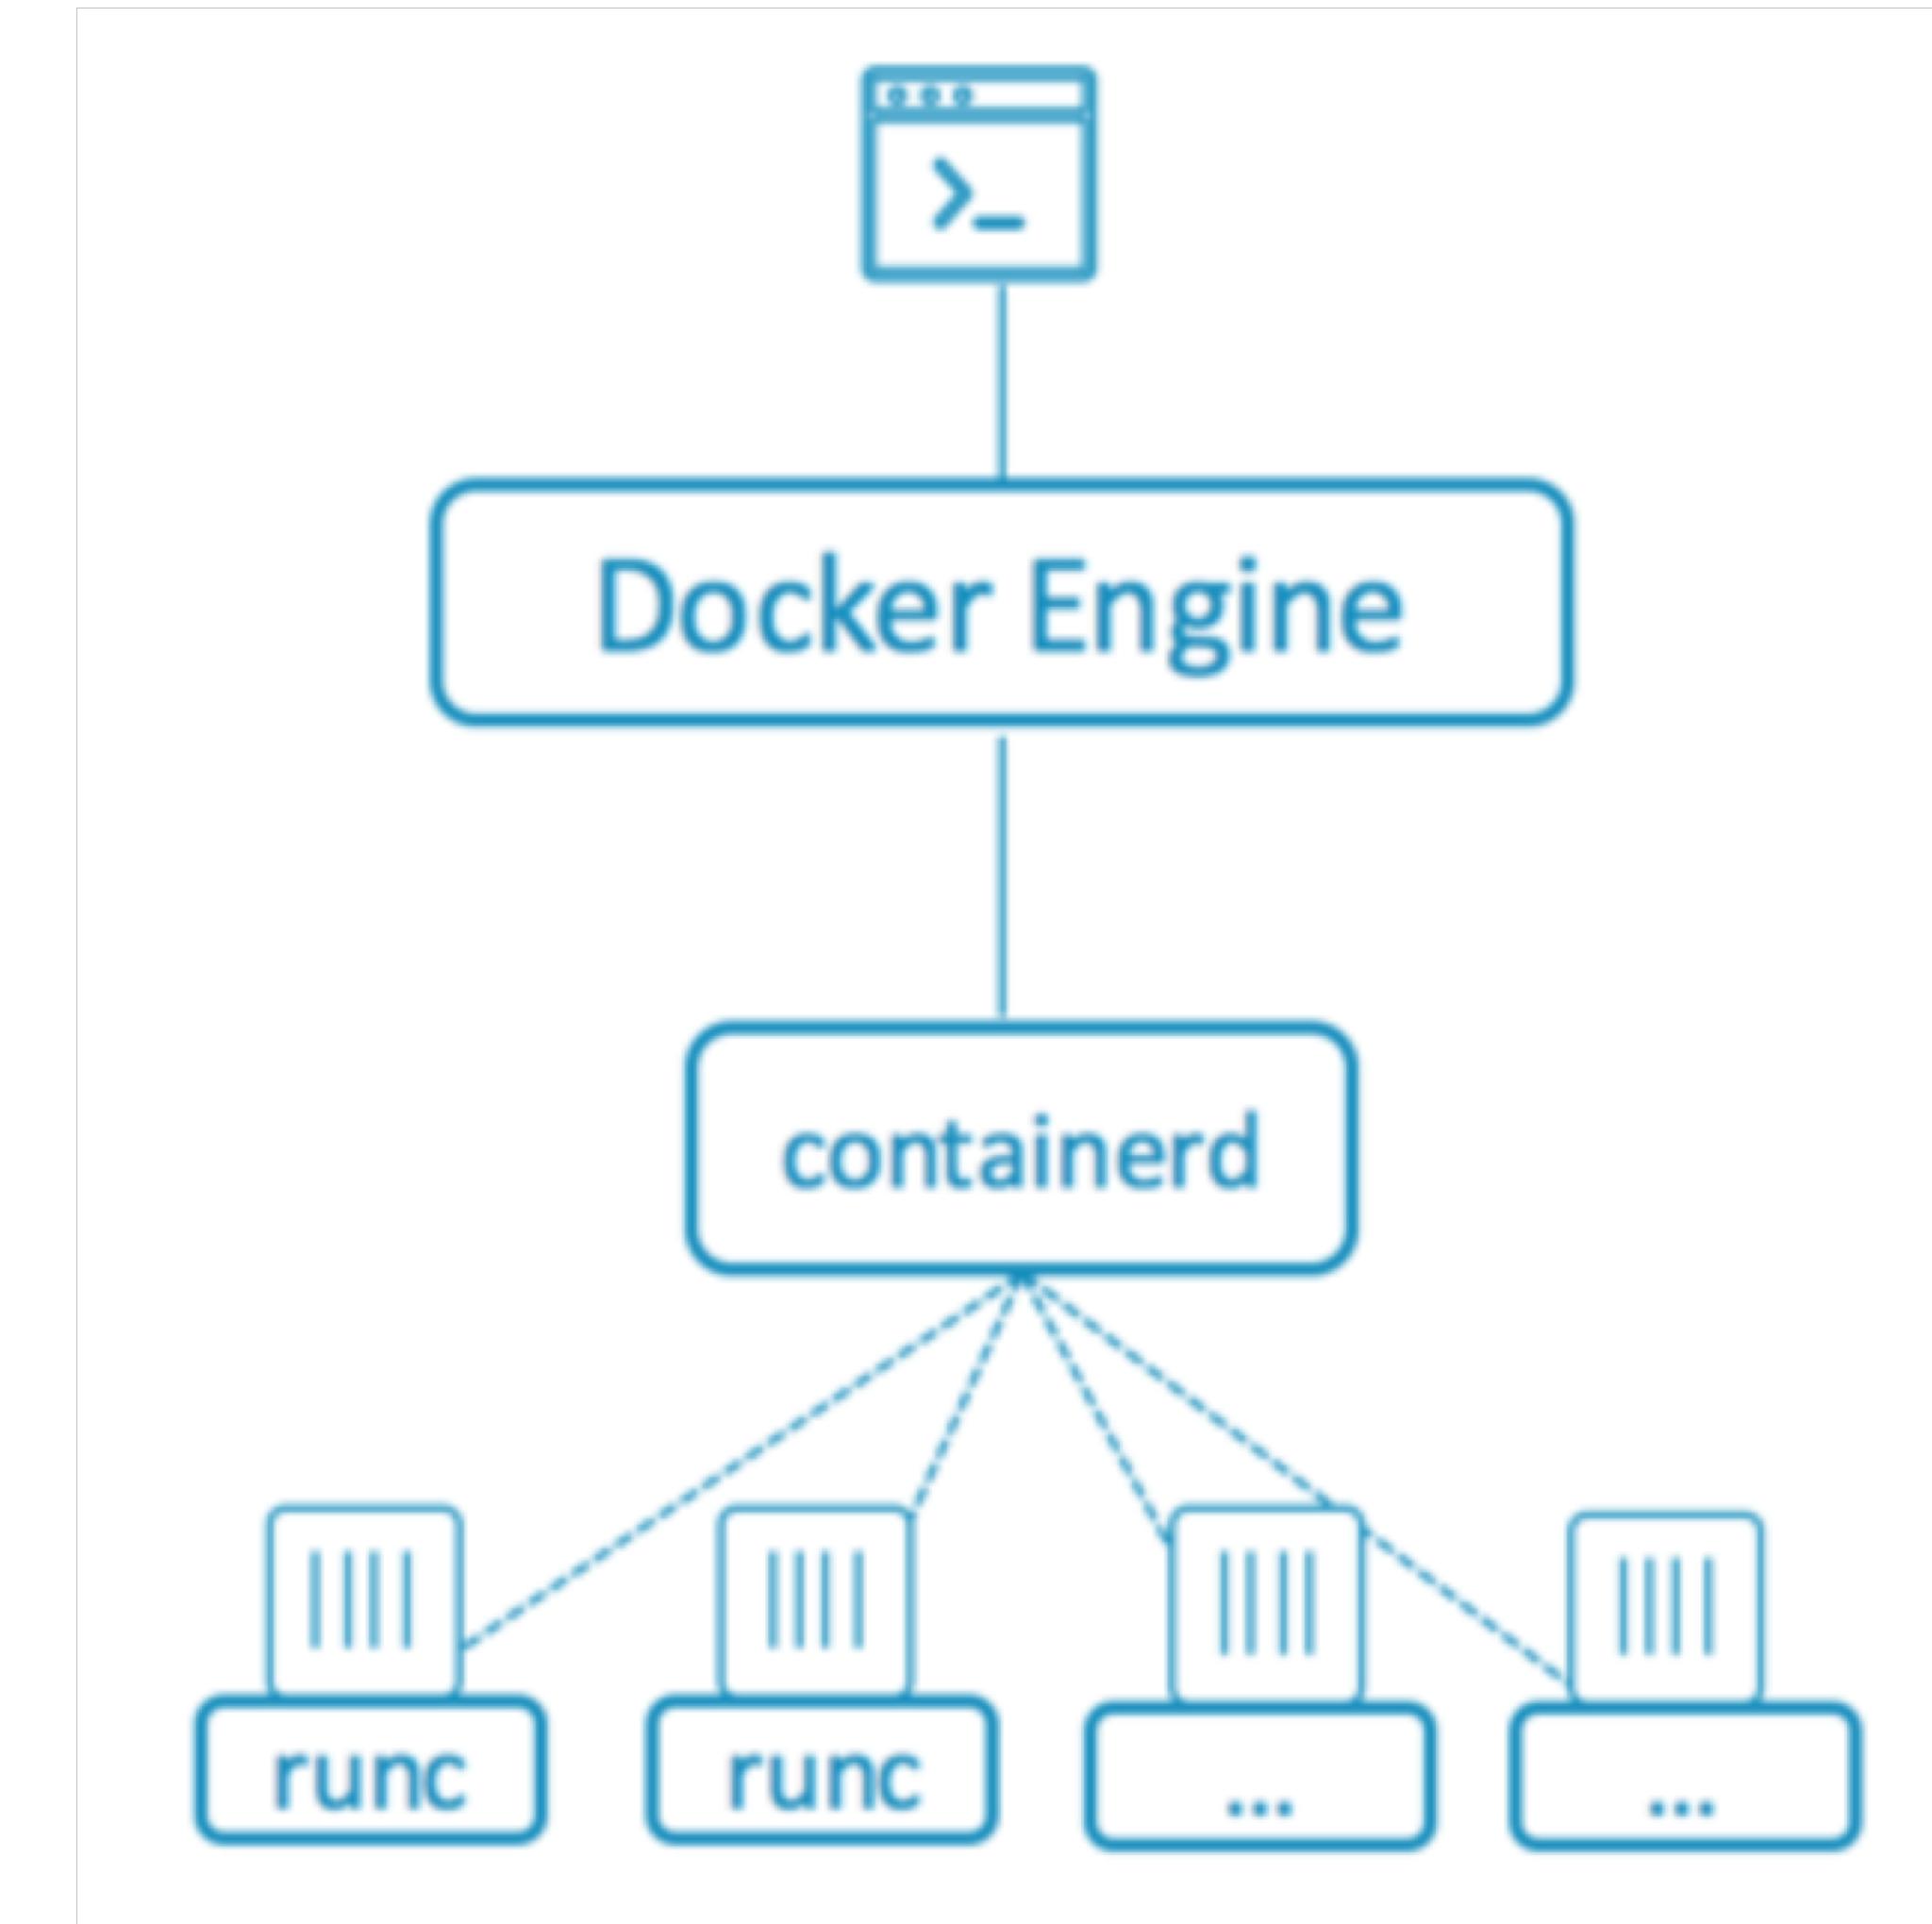


Runc

11

Runc is a lightweight, portable container runtime. It provides full support for Linux namespaces as well as native support for all security features available on Linux, such as SELinux, AppArmor, seccomp, and cgroups.

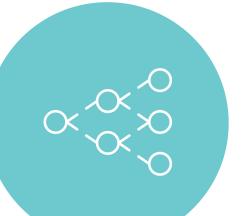
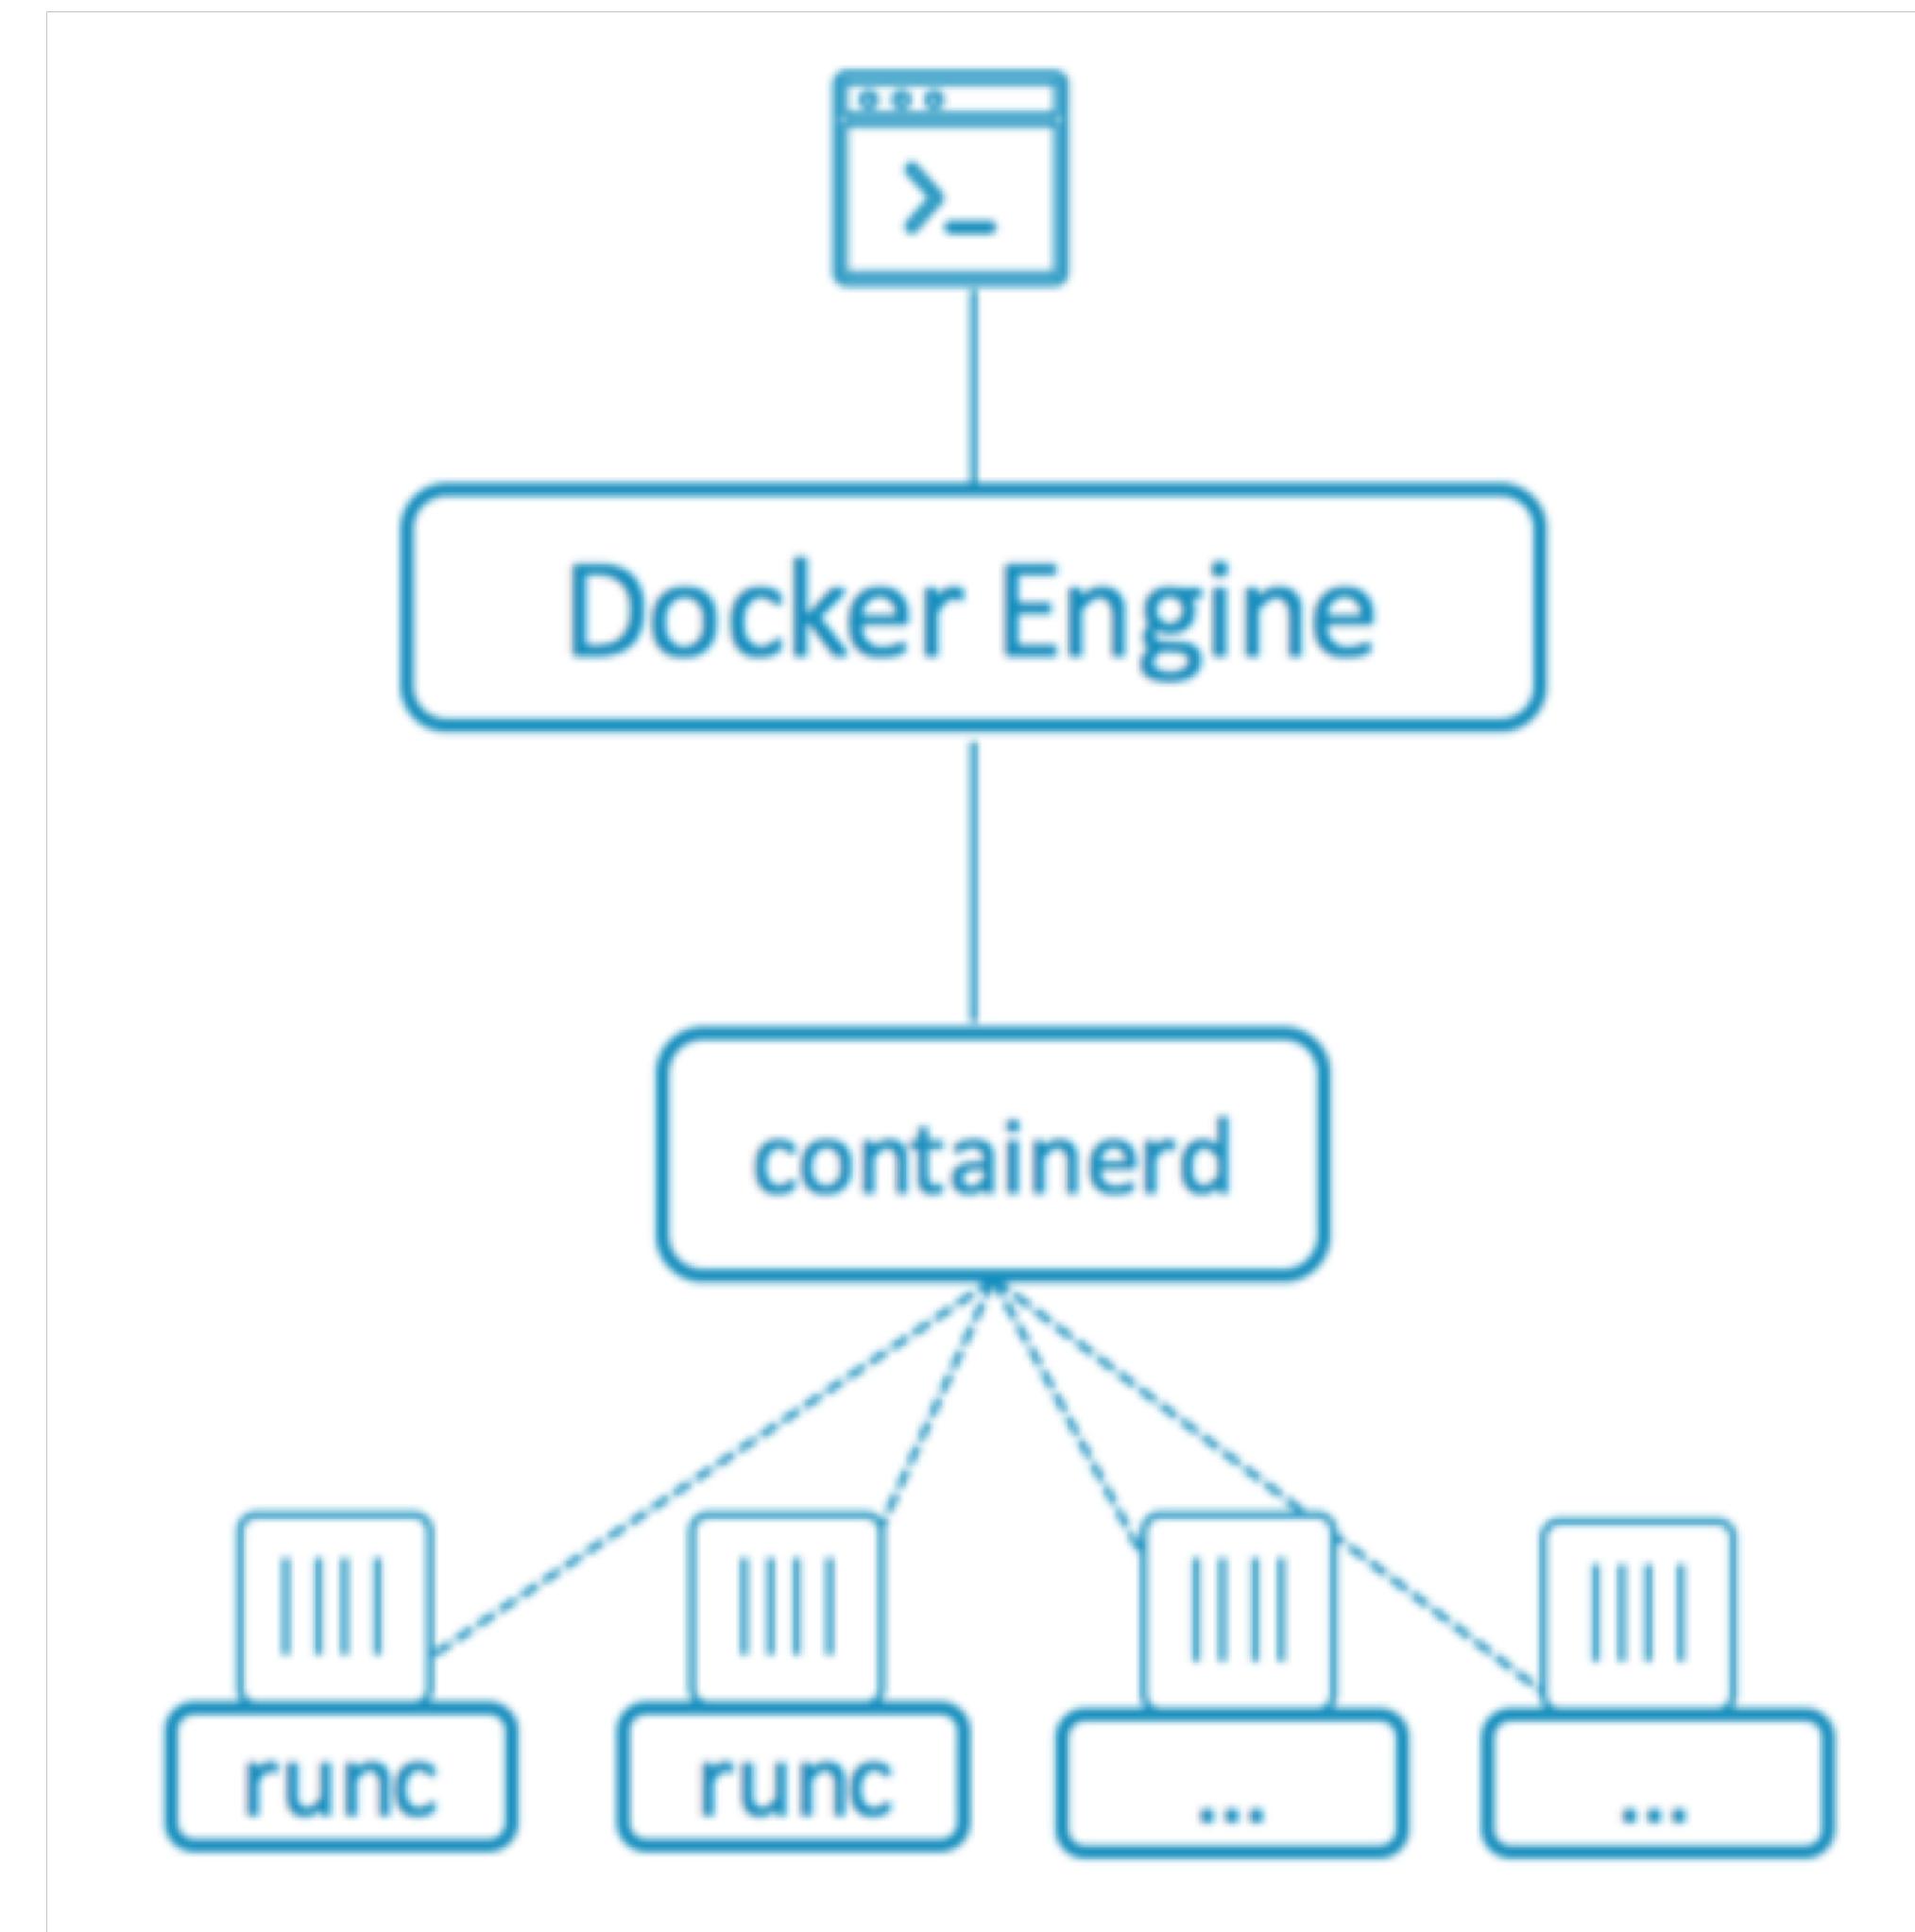
Runc is a tool for spawning and running containers according to the [Open Container Initiative\(OCI\)](#) specification. It is a formally specified configuration format, governed by the [Open Container Project \(OCP\)](#) under the auspices of the Linux Foundation.



Container.d

12

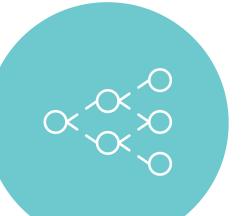
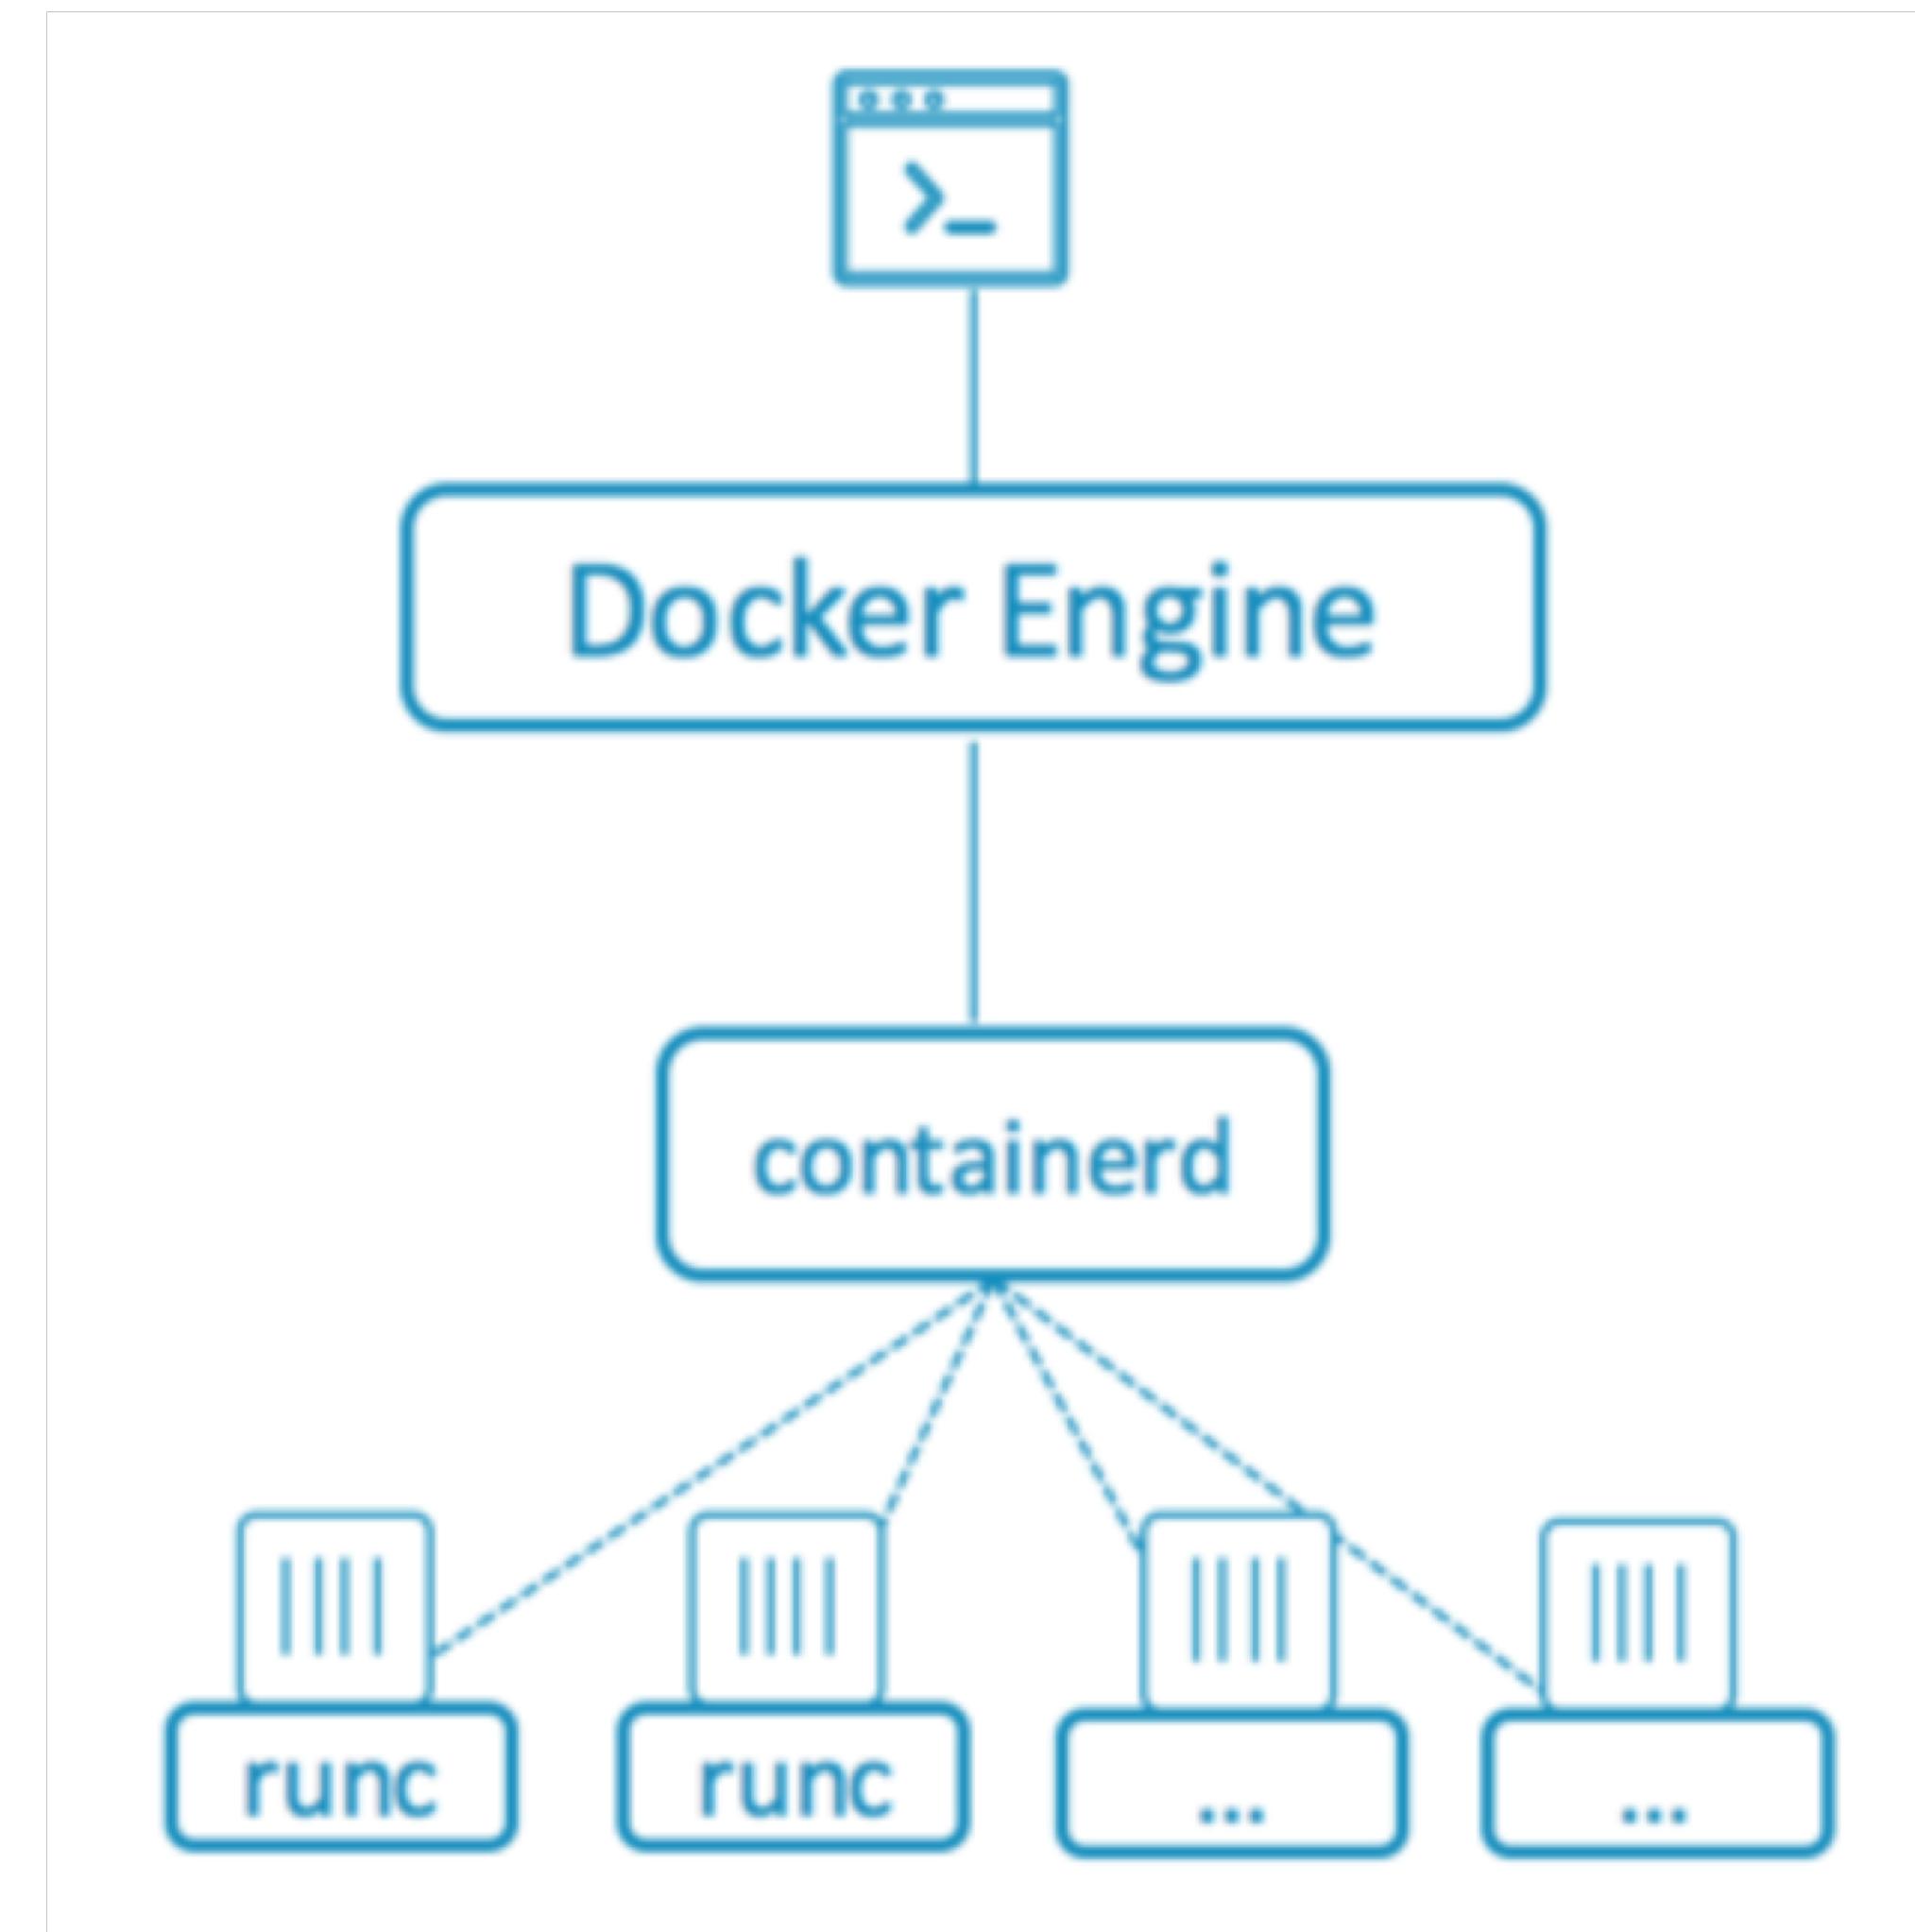
containerd builds on top of runc (container runtime), and adds higher-level features, such as image transfer and storage, container execution, and supervision, as well as network and storage attachments. With this, it manages the complete life cycle of containers. Containerd is the reference implementation of the OCI specifications and is by far the most popular and widely-used container runtime.



Container.d

13

containerd builds on top of runc (container runtime), and adds higher-level features, such as image transfer and storage, container execution, and supervision, as well as network and storage attachments. With this, it manages the complete life cycle of containers. Containerd is the reference implementation of the OCI specifications and is by far the most popular and widely-used container runtime.



Container – Commands -RUN

```
# Usage: docker container run [OPTIONS] IMAGE [COMMAND] [ARG...]
```

```
docker container run hello-world  
Unable to find image 'hello-world:latest' locally  
latest: Pulling from library/hello-world  
1b930d010525: Pull complete
```

Digest:

```
sha256:6540fc08ee6e6b7b63468dc3317e3303aae178cb8a45ed3123180328  
bcc1d20f
```

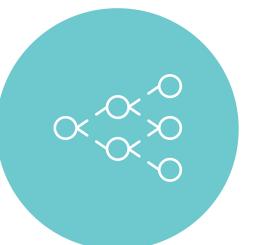
Status: Downloaded newer image for hello-world:latest

Docker unsuccessfully searched the local cache and then downloaded the hello-world:latest container image from the repository. Each container image description is made up of three parts:

- Docker registry host name
- Slash-separated name
- Tag name

Example hello-world image is stored this way – docker.io/hello-world:latest

Where docker.io is registry , hellow-world is name and latest is tag



Container – Commands -RUN

15

Usage: docker container run [OPTIONS] IMAGE [COMMAND] [ARG...]

To run container as detached

→ ~ docker container run -d hello-world

To run a container with name

→ ~ docker container run --name helloworld -d hello-world

Run a container with port mapping and forwarding : Port forwarding uses host: guest format

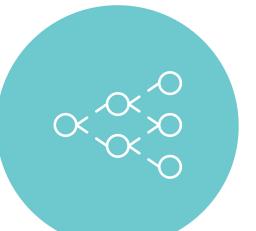
→ ~ docker container run --name helloworld -d -p 8080:80 hello-world

Running a command in the container during launch of container

→ ~ docker container run --name helloworld -d hello-world echo 'enjoy'

Running a container with interactive attached terminal

→ ~ docker container run --name helloworld -it hello-world



Exercise

16

- *1. Launch a ubuntu container and execute ls command in container*
- *2. Launch in centos container and ssh to this container*
- *3. Install git to this container*

Container – Commands - LS

17

Usage: docker container ls [OPTIONS]

Docker container ls will list all the container currently active

→ ~ docker container ls

CONTAINER					
ID	IMAGE	COMMAND	CREATE	NA	
D	STATUS	PORTS	NA		
MES					

Lists all container even the killed and exited ones

→ ~ docker container ls -a

List last 2 container

→ ~ docker container ls -n 2

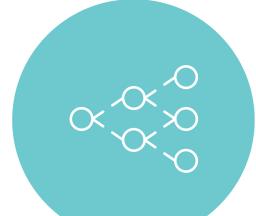
Just list only id of container

~ docker container ls -n 2 -q

Aliases:
ls, ps, list

Options:

- | | |
|--|--|
| -a, --all
running) | Show all containers (default shows just running) |
| -f, --filter filter
provided | Filter output based on conditions |
| --format string
template | Pretty-print containers using a Go template |
| -n, --last int
all states) (default -1) | Show n last created containers (includes all states) |
| -l, --latest
(includes all states) | Show the latest created container |
| --no-trunc | Don't truncate output |
| -q, --quiet | Only display numeric IDs |
| -s, --size | Display total file sizes |



Exercise

18

- *1. List all container created*
- *2.Just list the container which are running*
- *3. List containers with name helloworld and mysql*

Container – Commands - RM

19

Usage: docker container rm [OPTIONS]

Docker container rm, will remove the container from system

→ ~ docker container rm <container_ids>

Usage: docker container rm [OPTIONS] CONTAINER
[CONTAINER...]

Remove one or more containers

Options:

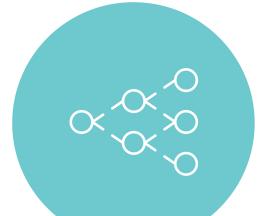
-f, --force Force the removal of a running container
(uses SIGKILL)

-l, --link Remove the specified link

-v, --volumes Remove the volumes associated with the
container

Removing the running container without stopping them

→ ~ docker container rm -f 98378377a



Exercise

20

- . *1. Remove all container running non –running*

Container – Commands – Create/Stop/Start

Usage: docker container create [OPTIONS] IMAGE

→ ~ docker container create --name nginx-test -p 8080:80 nginx

Container create command creates a container but it will not run. To run a container you need to use start command

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
78433ad748e2	nginx	"nginx -g 'daemon off;"	6 seconds ago	Created		

To run a container

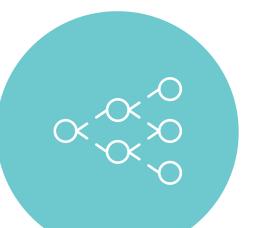
→ ~ docker container start **78433ad748e2**

To stop a container

→ ~ docker container stop 78433ad748e2

To kill a container

→ ~ docker container kill 78433ad748e2



Exercise

22

- *1. Run mongoDB, create a collection in MongoDB. Stop the container and start it again, verify collection is still there.*
- *Create mysql container and run and kill it.*

Container – Commands – LOGS

```
# Usage: docker container logs [OPTIONS] CONTAINER
```

```
→ ~ docker container logs --follow --timestamps 78433ad748e2
```

```
→ ~ docker container logs --follow --timestamps 78433ad748e2
2019-07-24T18:00:42.225431266Z 172.17.0.1 - - [24/Jul/2019:18:00:42 +0000] "GET / HTTP/1.1" 200
612 "-" "curl/7.54.0" "-"
```

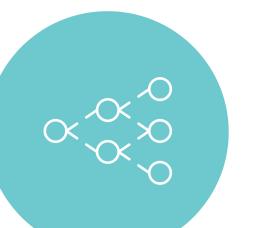
```
# the short form of the command
```

```
docker container logs -f -t web-server
```

```
# the short form of the command
```

```
→ ~ docker container logs --follow --tail 5 78433ad748e2
```

The --details, --follow, --timestamps, and --tail parameters are all optional,



Exercise

24

- *1. Run mongoDB, get the logs for the container*

Container – Commands – TOP

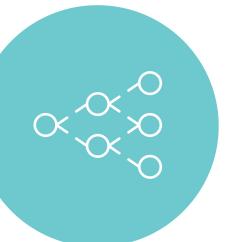
25

```
# using the new syntax  
# Usage: docker container top CONTAINER [ps OPTIONS]  
docker container top web-server
```

```
# using the old syntax  
docker top web-server
```

List processes are running inside a container

```
→ ~ docker container top 78433ad748e2  
PID          USER          TIME        COMMAND  
2673          root         0:00      nginx: master process nginx -g daemon off;  
2708          101          0:00      nginx: worker process
```



Inspecting a container

```
# using the new syntax
# Usage: docker container inspect [OPTIONS] CONTAINER [CONTAINER...]
```

→ ~ docker container inspect 78433ad748e2

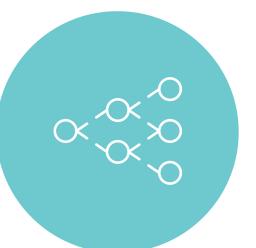
Inspect will dump a json with all details.

Get details of particular json node and format

```
→ ~ docker container inspect --format '{{json .State}}' 78433ad748e2 | jq
{
  "Status": "running",
  "Running": true,
  "Paused": false,
  "Restarting": false,
  "OOMKilled": false,
  "Dead": false,
  "Pid": 2673,
  "ExitCode": 0,
  "Error": "",
  "StartedAt": "2019-07-24T17:59:03.780192719Z",
  "FinishedAt": "0001-01-01T00:00:00Z"
}
```

Get particular value

→ ~ docker container inspect --format '{{json .State}}' 78433ad748e2 | jq '.StartedAt'
 "2019-07-24T17:59:03.780192719Z"



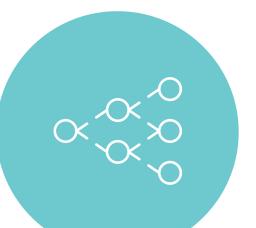
Inspecting a container

Get IP address of container using inspect

```
→ ~ docker container inspect -f '{{json .NetworkSettings}}'  
78433ad748e2 | jq .IPAddress  
"172.17.0.2"  
→ ~
```

Get ports of container

```
→ ~ docker container inspect --format '{{json .NetworkSettings}}'  
78433ad748e2 | jq '.Ports'  
{  
  "80/tcp": [  
    {  
      "HostIp": "0.0.0.0",  
      "HostPort": "8080"  
    }  
  ]  
}
```

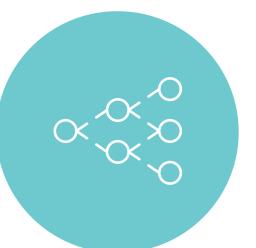


Getting a container stats

Provides live, continually-updated usage statistics for one or more running containers

→ ~ docker container stats 78433ad748e2

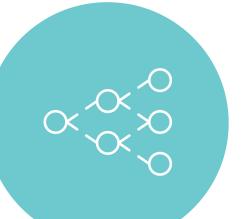
CONTAINER ID	NAME	CPU %	MEM USAGE / LIMIT	MEM %	NET
I/O	BLOCK I/O	PIDS			
78433ad748e2	nginx-test	0.00%	1.895MiB / 1.952GiB	0.09%	2.02kB /
1.27kB	0B / 0B	2			



Exercise

29

- *1. Run mongo container.*
- *2. Run a mysql container*
- *3. list ip and port for the containers*
- *4. get CPU and % Memory used by them.*



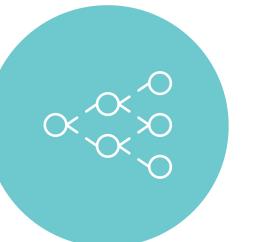
Attaching a container

```
# run a container detached  
→ ~ docker container run -idt --name nginx-test3 -p 8081:80 nginx  
50908027af2d234d6b5c374e8a746044f866e3fbe6f7b6097a99909ce60eae6b
```

Attach container to the terminal

```
→ ~ docker container attach 50908027
```

```
# issue a Ctrl + PQ keystroke to detach
```



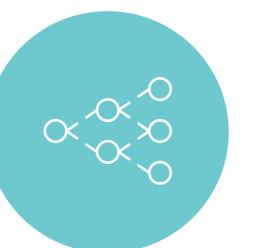
Docker container exec

Sometimes, when you have a container running detached, you might want to get access to it, but don't want to attach to the executing command. You can accomplish this by using the container exec command

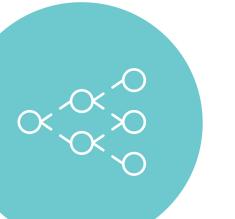
```
→ ~ docker container exec 78433ad748e2 ls
```

Run exec with interactive mode, example running shell

```
→ ~ docker container exec -it 78433ad748e2 /bin/bash  
root@78433ad748e2:/#
```

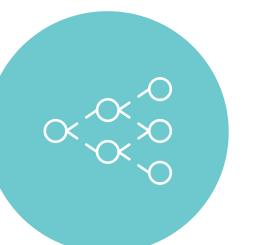


Docker Images



Images

In Linux, everything is a file. The whole operating system is basically a filesystem with files and folders stored on the local disk. This is an important fact to remember when looking at what container images are. As we will see, an image is basically a big tarball containing a filesystem. More specifically, it contains a layered filesystem



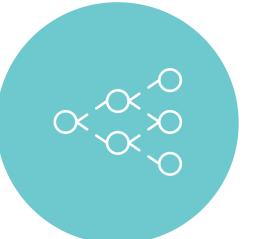
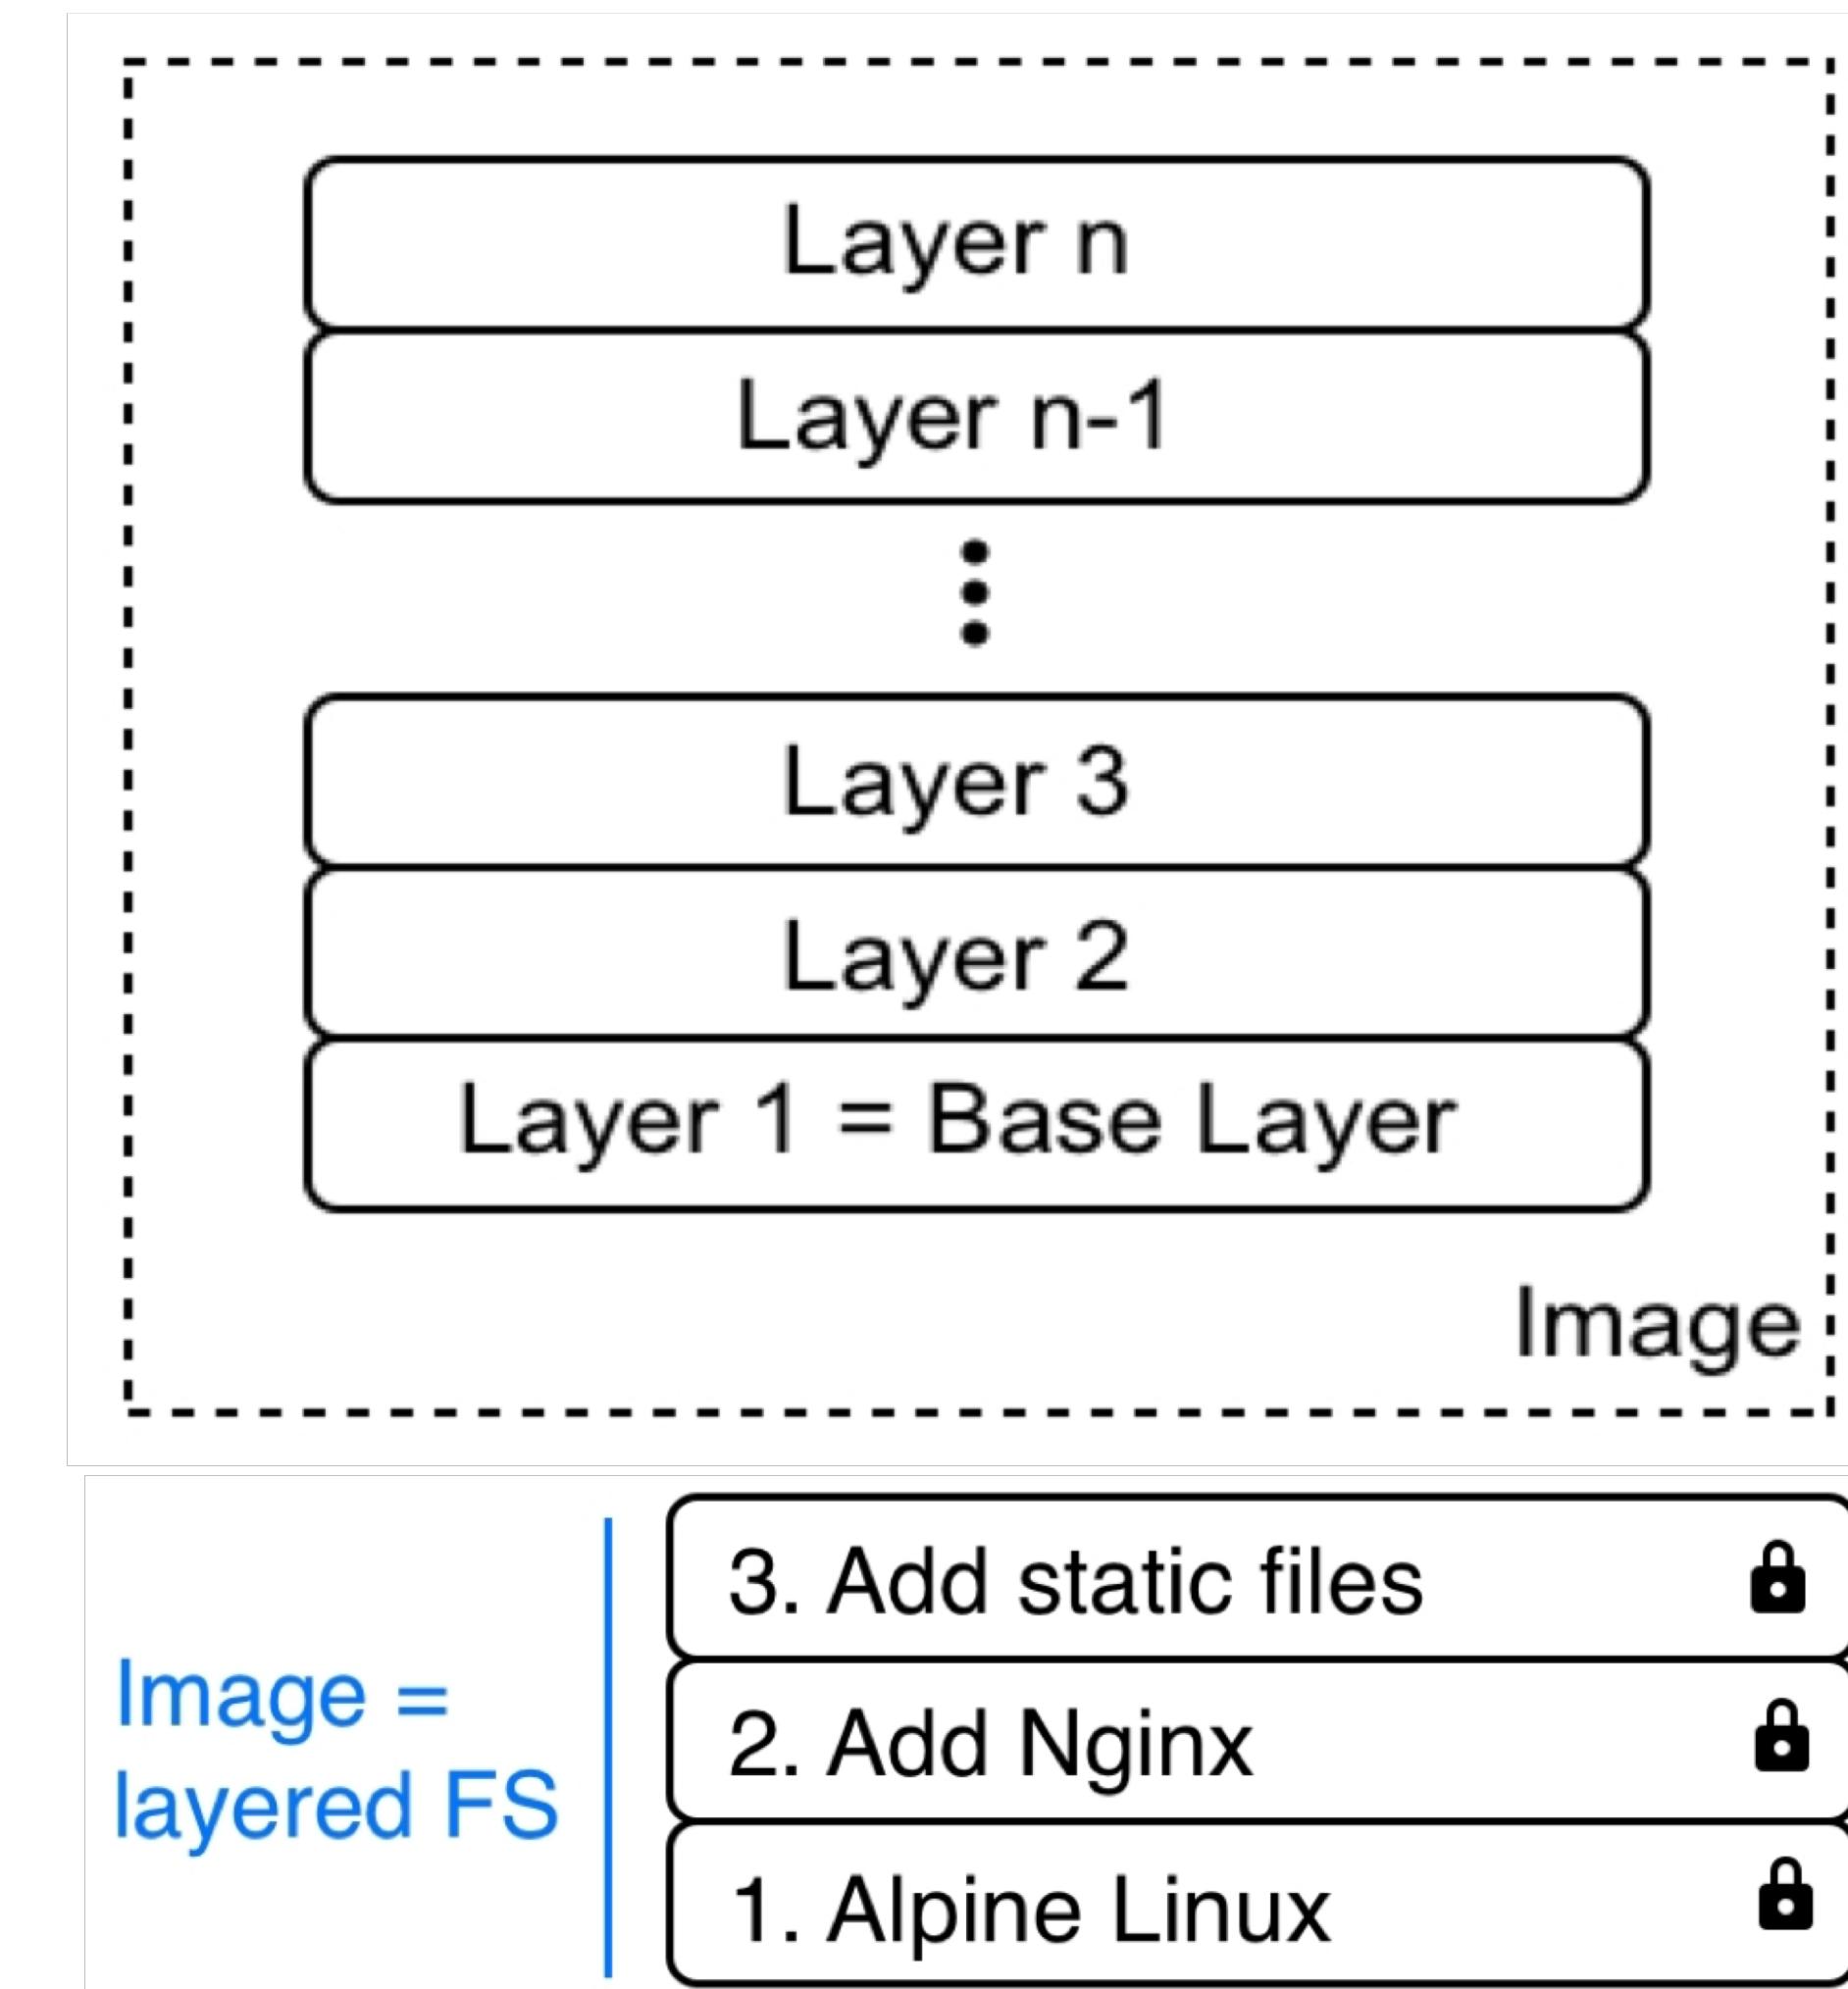
Layered FileSystem

Container images are templates from which containers are created. These images are not just one monolithic block, but are composed of many layers. The first layer in the image is also called the base layer:

Each individual layer contains files and folders.

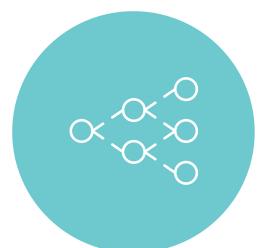
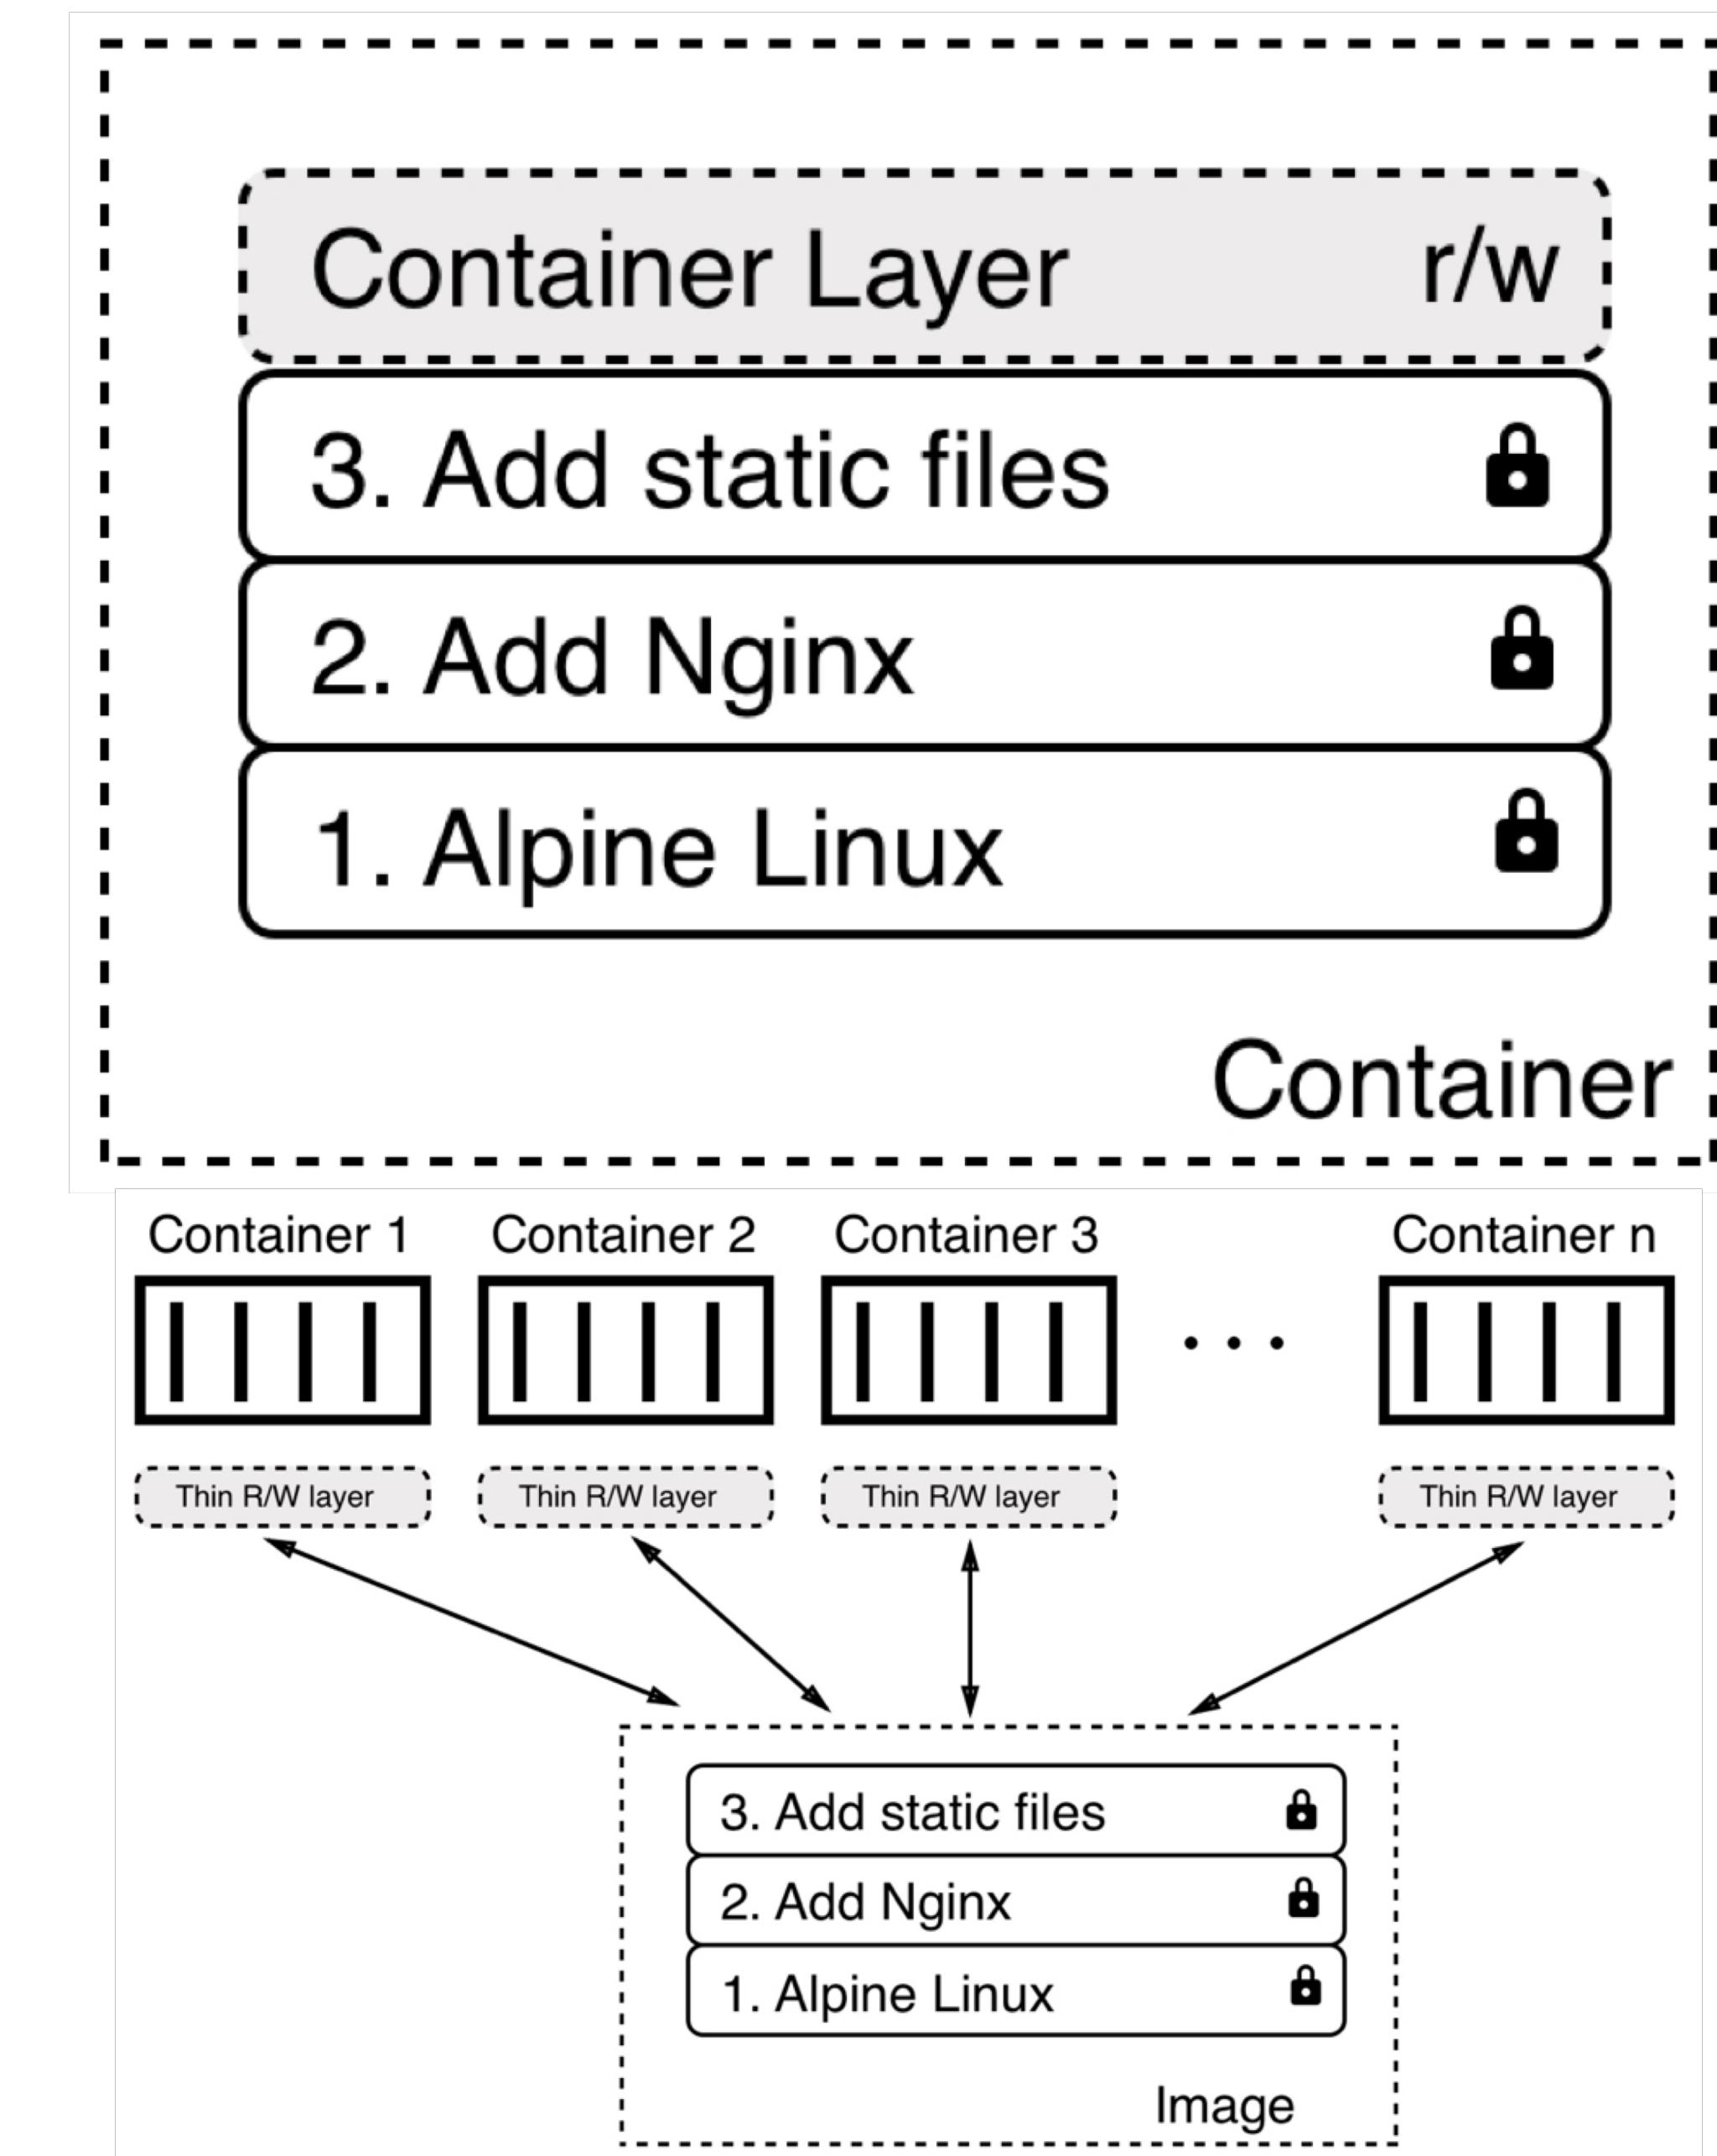
Each layer only contains the changes to the filesystem with respect to the underlying layers

In the following image, we can see what a custom image for a web application using Nginx as a web server could look like:



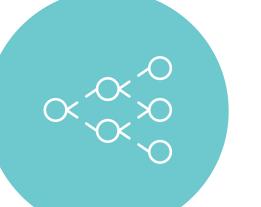
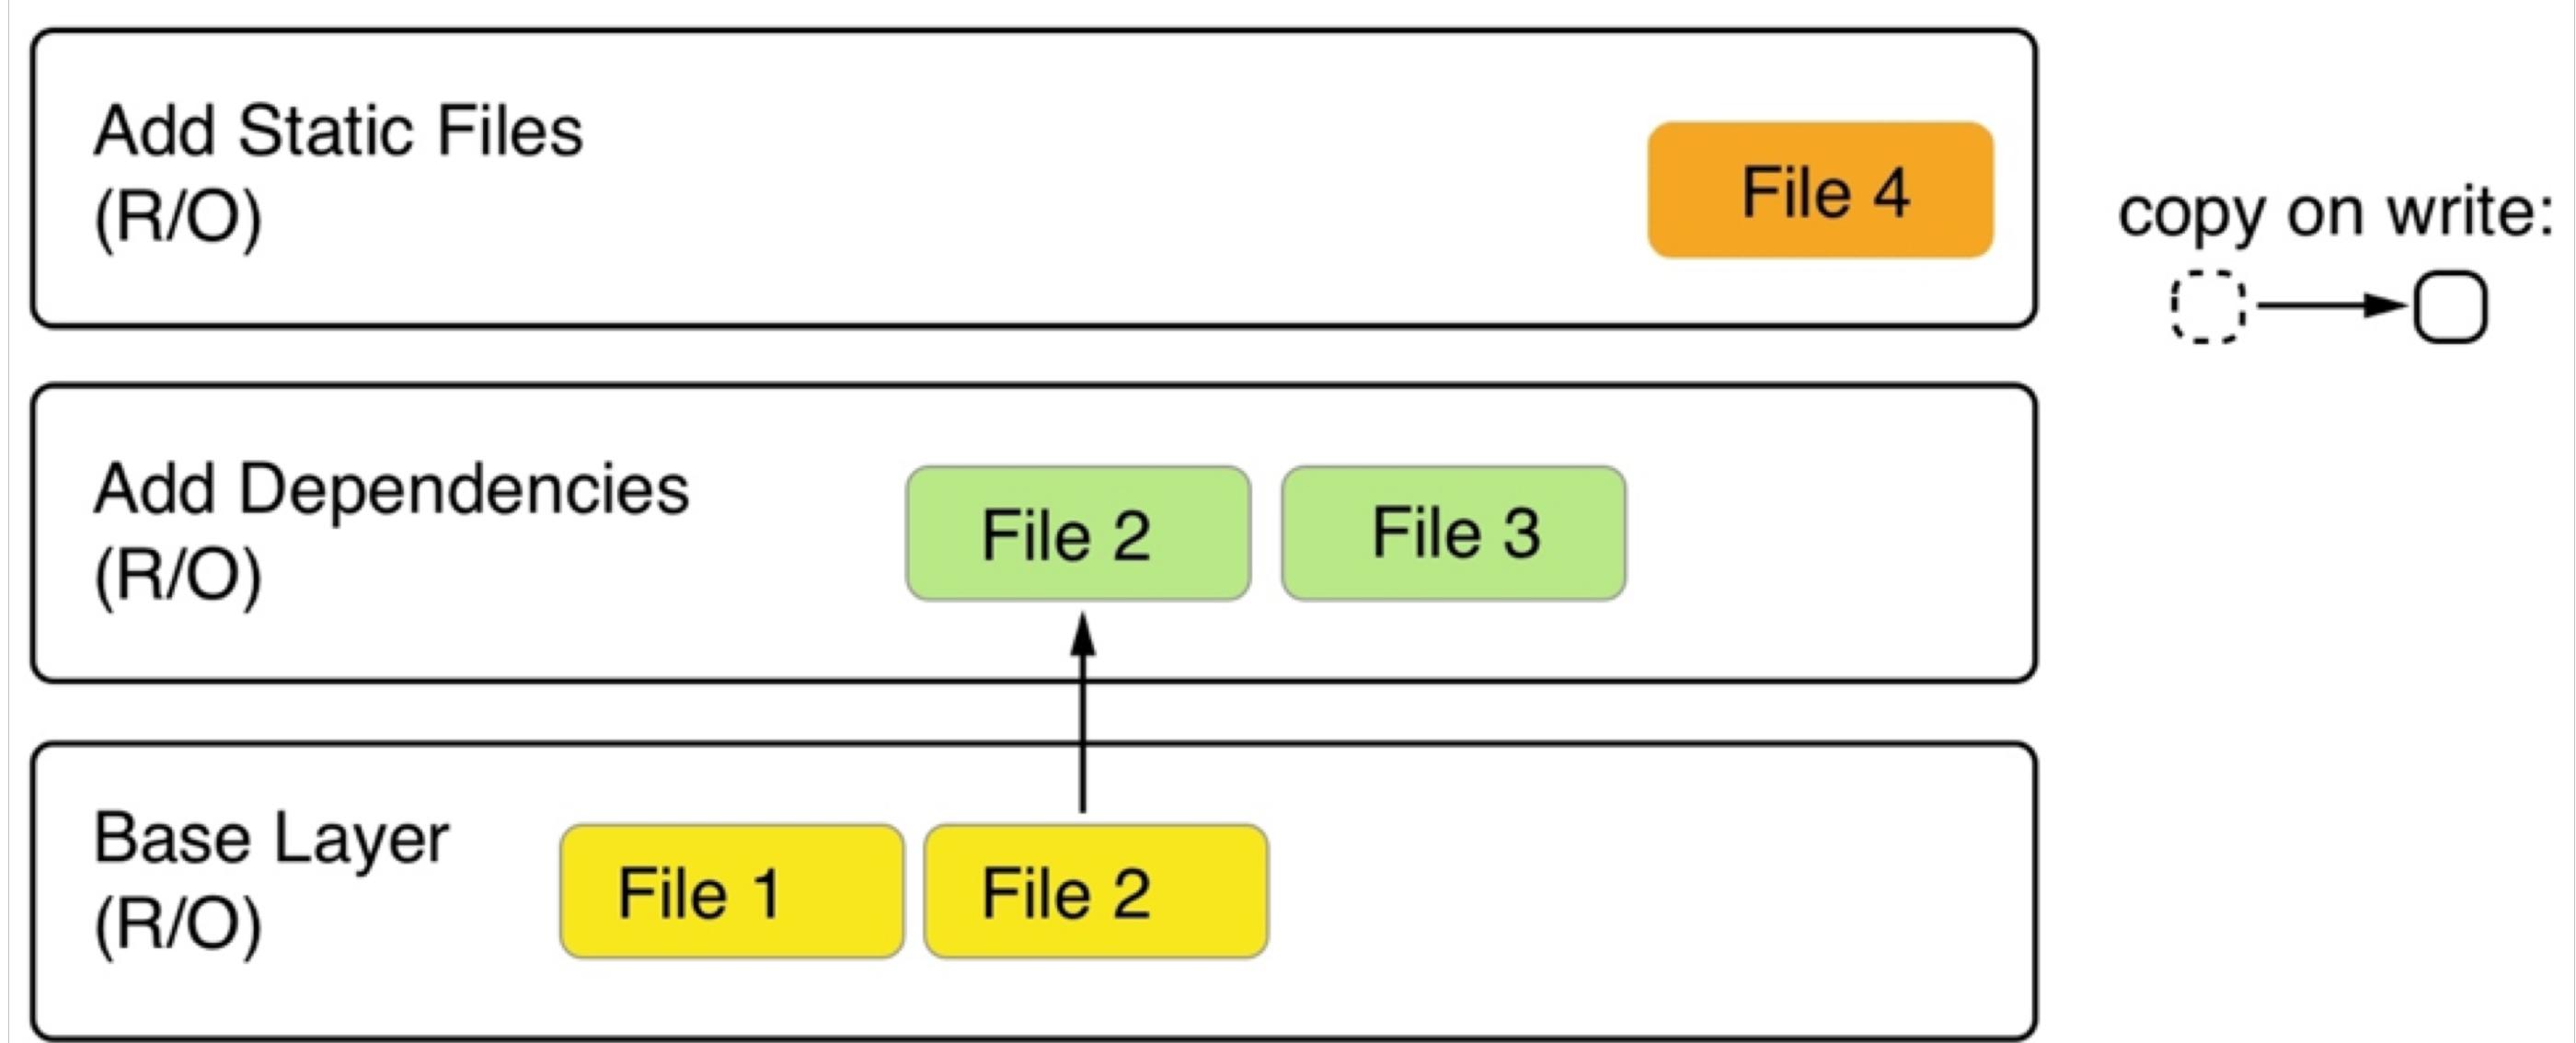
Writable Container Layer

container image is made of a stack of immutable or read-only layers. When the Docker engine creates a container from such an image, it adds a writable container layer on top of this stack of immutable layers. Our stack now looks as follows



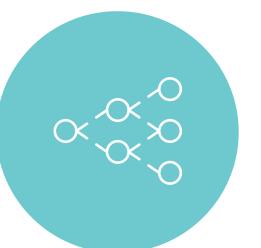
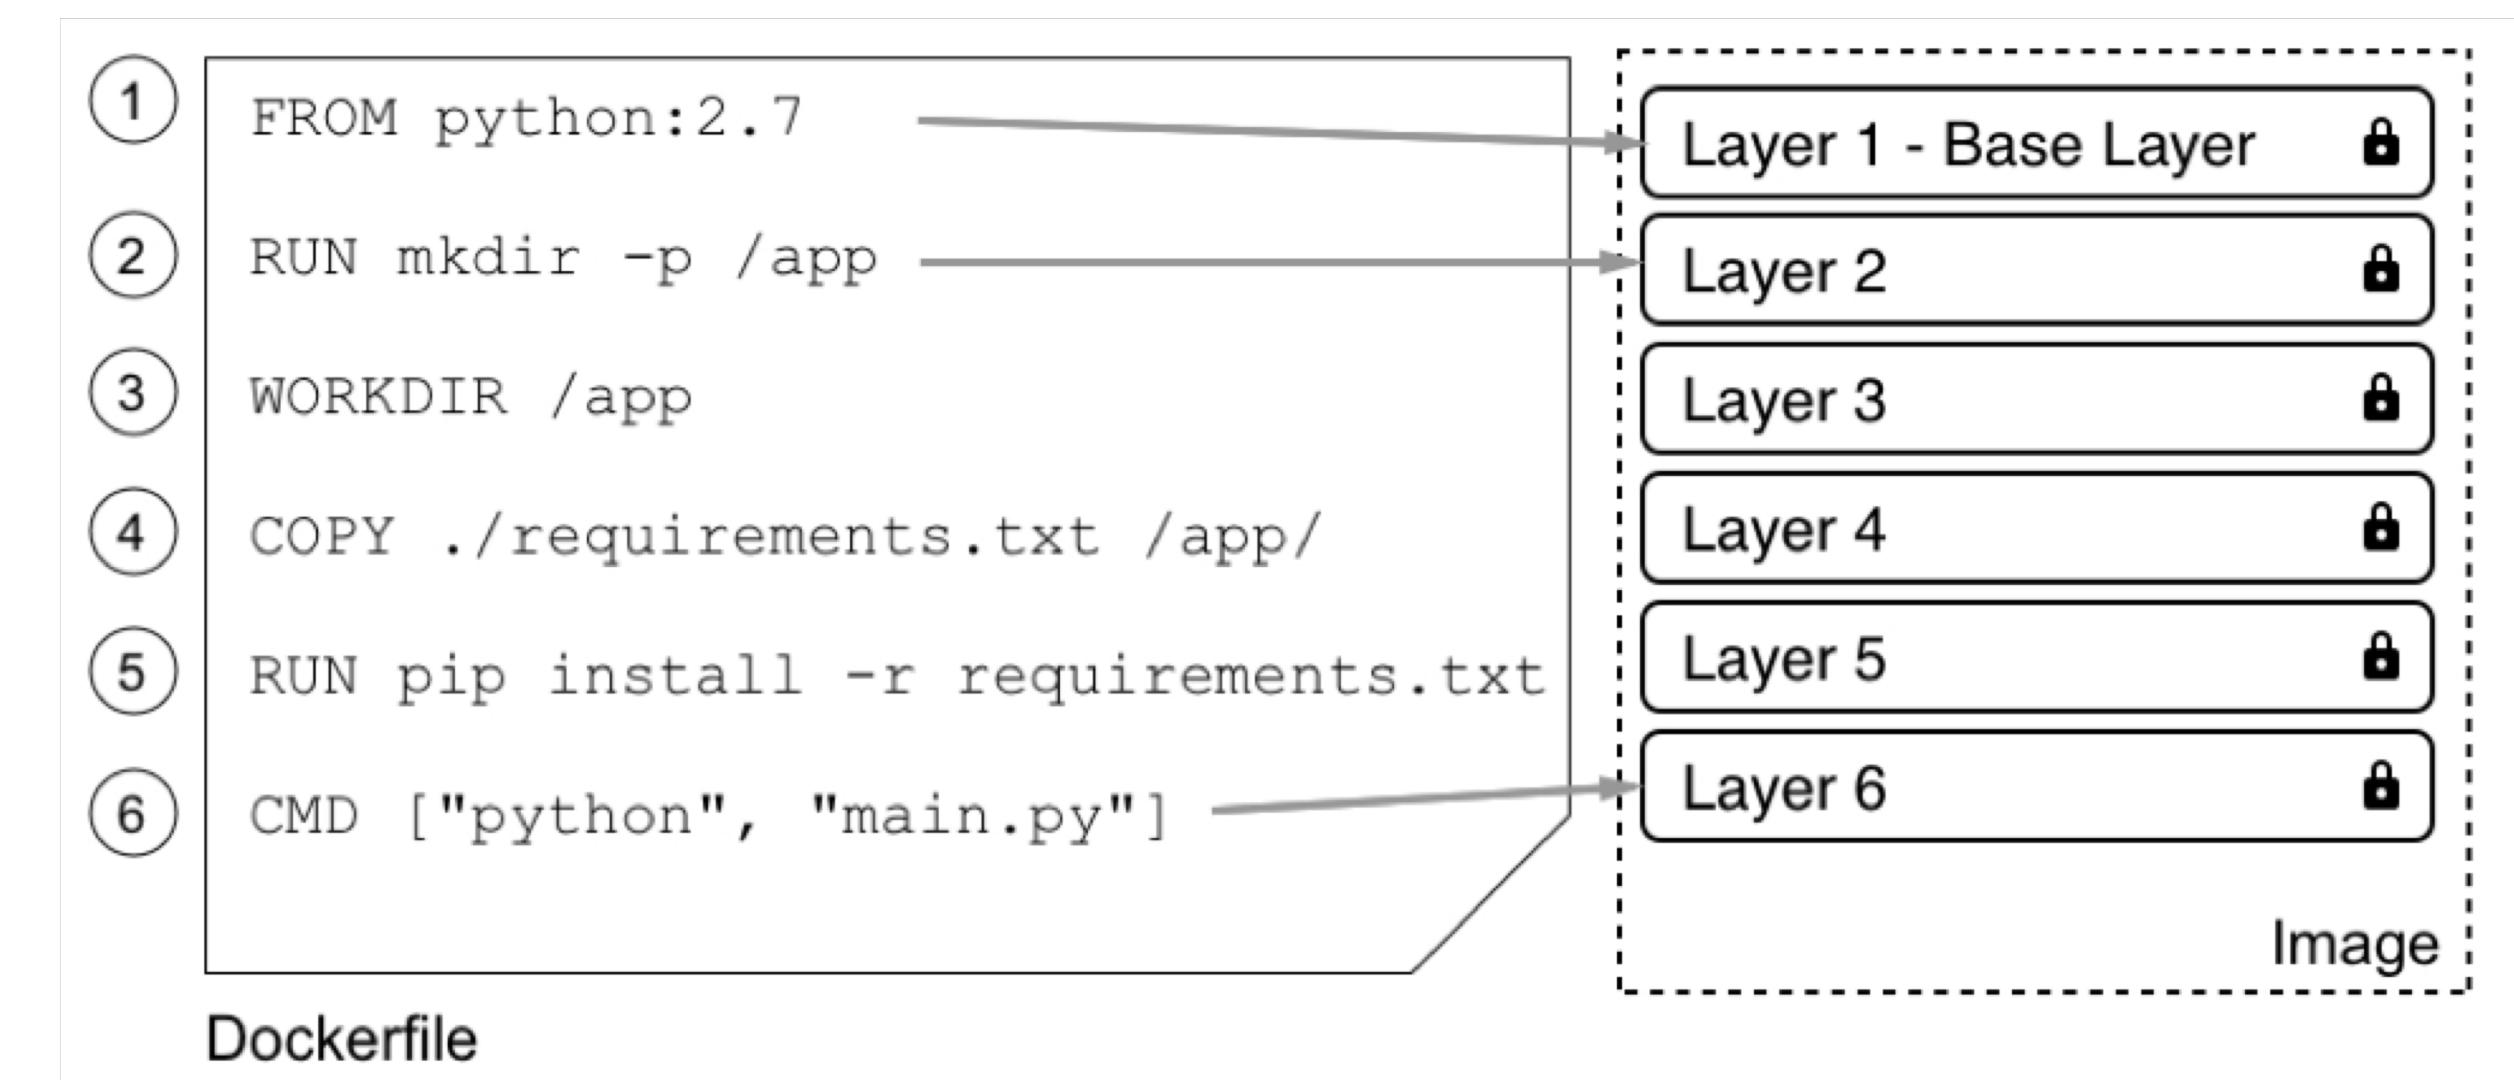
Copy on Write

Docker uses the copy-on-write technique when dealing with images. Copy-on-write is a strategy of sharing and copying files for maximum efficiency. If a layer uses a file or folder that is available in one of the low-lying layers, then it just uses it. If, on the other hand, a layer wants to modify, say, a file from a low-lying layer, then it first copies this file up to the target layer and then modifies it. In the following figure, we can see a glimpse of what this means:



Dockerfile

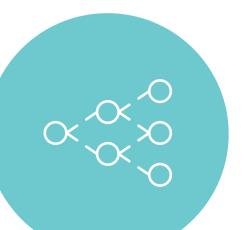
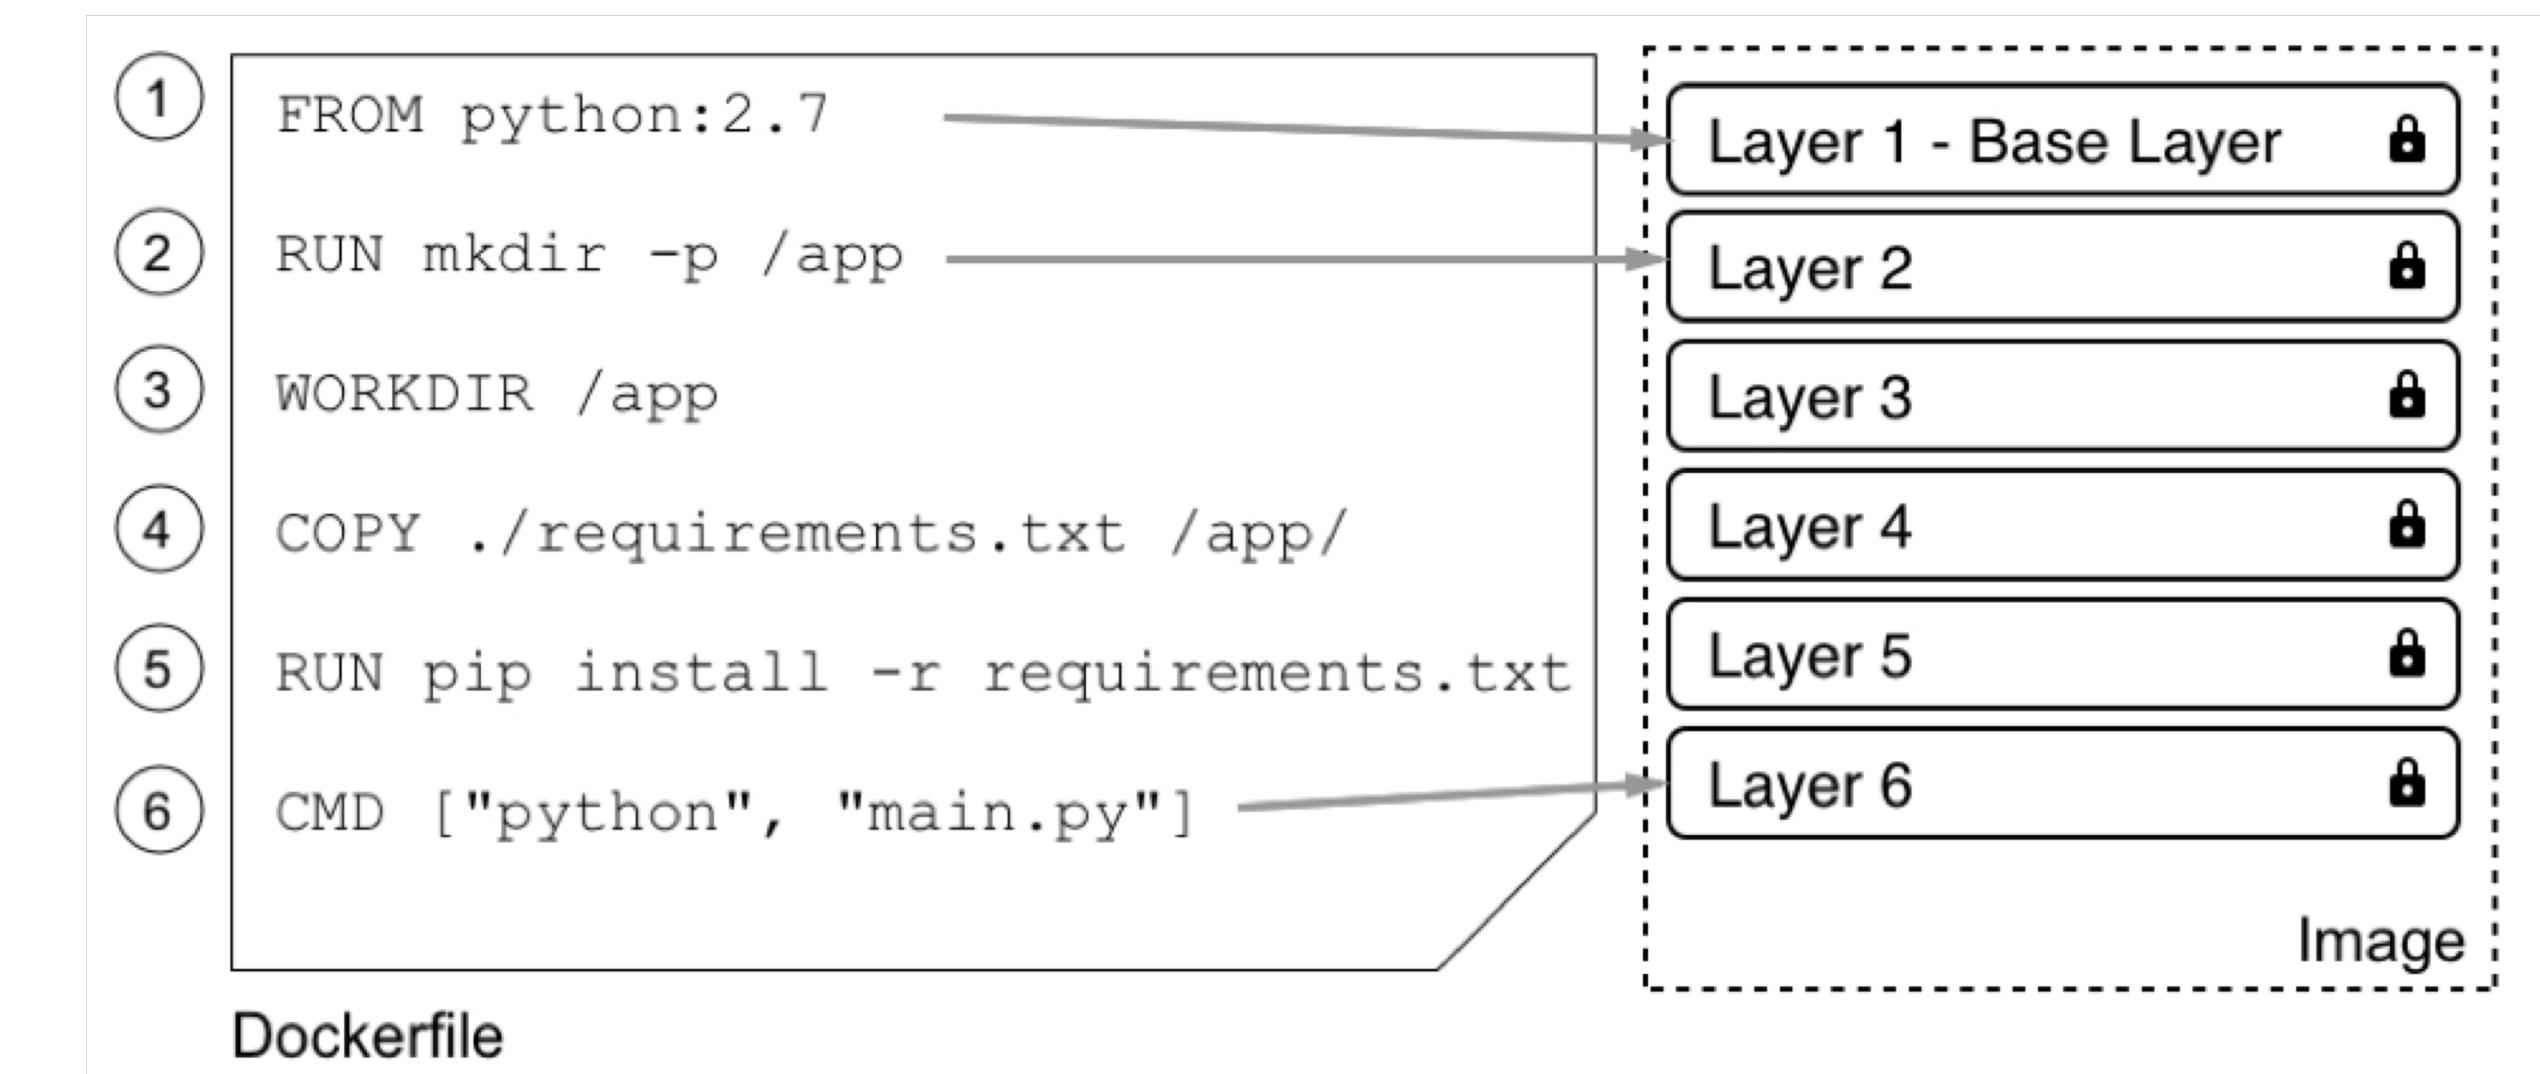
A Dockerfile is a text file that contains instructions used by the Docker daemon to create a Docker image. The instructions are defined using a type of value pair syntax. Each one has an instruction word followed by the parameters for that instruction. Every command gets its own line in the Dockerfile. Although the Dockerfile instructions are not case-sensitive, there is a well-used convention that the instruction word is always uppercase.



FROM Instruction

Every Dockerfile must have a FROM instruction, and it must be the first instruction in the file.

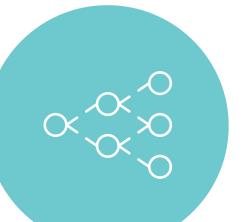
The FROM instruction sets the base for the image being created and instructs the Docker daemon that the base of the new image should be the existing Docker image specified as the parameter



The LABEL instruction

The LABEL instruction is a way to add metadata to your Docker image. This instruction adds embedded key-value pairs to the image. The LABEL instruction adds a zero-byte-sized layer to the image when it is created. An image can have more than one LABEL, and each LABEL instruction can provide one or more LABELs. The most common use for the LABEL instruction is to provide information about the image maintainer.

```
LABEL maintainer="Earl Waud  
<earlwaud@mycompany.com>"  
LABEL "description"="My development Ubuntu  
image"  
LABEL version="1.0"  
LABEL label1="value1" \  
label2="value2" \  
label3="value3"  
LABEL my-multi-line-label="Labels can span \  
more than one line in a Dockerfile."  
LABEL support-email="support@mycompany.com"  
support-phone="(123) 456-7890"
```



The COPY instruction

The COPY instruction is used to copy files and folders into the Docker image being built

Without the --chown parameter, the owner ID and group ID will both be set to 0.

The <src> or source is a filename or folder path and is interpreted to be relative to the context of the build. We will talk more about the build context later in this chapter, but for now, think of it as where the build command is run. The source may include wildcards

The <dest> or destination is a filename or path inside of the image being created

```
# COPY instruction syntax
COPY [--chown=<user>:<group>] <src>... <dest>
# Use double quotes for paths containing whitespace)
COPY [--chown=<user>:<group>] [<src>,... <dest>]
```

```
# COPY instruction Dockerfile for Docker Quick Start
FROM alpine:latest
LABEL maintainer="Earl Waud <earlwaud@mycompany.com>"
LABEL version=1.0
# copy multiple files, creating the path
#/theqsg/files" in the process
COPY file* theqsg/files/
# copy all of the contents of folder "folder1" to
#/theqsg/
# (but not the folder "folder1" itself)
COPY folder1 theqsg/
# change the current working directory in the image to
#/theqsg"
WORKDIR theqsg
# copy the file special1 into "/theqsg/special-files/"
COPY --chown=35:35 special1 special-files/
# return the current working directory to "/"
WORKDIR /
CMD ["sh"]
```

The ADD instruction

The ADD instruction is used to copy files and folders into the Docker image being built

ADD instruction can actually do more than the COPY instruction. The more is dependent upon the values used for the source input. With the COPY instruction, the source can be files or folders. However, with the ADD instruction, the source can be files, folders, a local .tar file, or a URL.

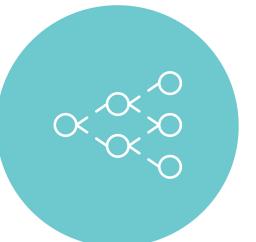
ADD instruction syntax

```
ADD [--chown=<user>:<group>] <src>... <dest>
```

Use double quotes for paths containing whitespace)

```
ADD [--chown=<user>:<group>] ["<src>",... "<dest>"]
```

```
# ADD instruction Dockerfile for Docker Quick Start
FROM alpine
LABEL maintainer= "Atin<atin@pragra.co>"
LABEL version=3.0
ADD https://github.com/docker-library/hello-
world/raw/master/amd64/hello-world/hello /
RUN chmod +x /hello
CMD ["/hello"]
```

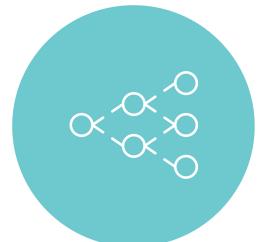


The ENV instruction

ENV instruction is used to define environment variables that will be set in the running containers created from the image being built. The variables are defined using typical key-value pairs. A Dockerfile can have one or more ENV instructions. Here is the ENV instruction syntax:

```
# ENV instruction syntax
ENV <key> <value>
ENV <key>=<value> ..
```

```
# ENV instruction Dockerfile for Docker Quick Start
FROM alpine
LABEL maintainer="Atin <atin@pragra.co>"
ENV appDescription This app is a sample of using ENV instructions
ENV appName=env-demo
ENV note1="The First Note First" note2=The\ Second\ Note\ Second \
note3="The Third Note Third"
ENV changeMe="Old Value"
CMD ["sh"]
```

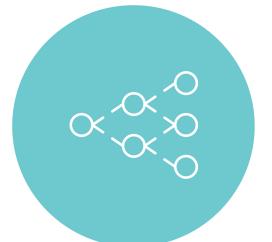


The ENV instruction

ENV instruction is used to define environment variables that will be set in the running containers created from the image being built. The variables are defined using typical key-value pairs. A Dockerfile can have one or more ENV instructions. Here is the ENV instruction syntax:

```
# ENV instruction syntax
ENV <key> <value>
ENV <key>=<value> ..
```

```
# ENV instruction Dockerfile for Docker Quick Start
FROM alpine
LABEL maintainer="Atin <atin@pragra.co>"
ENV appDescription This app is a sample of using ENV instructions
ENV appName=env-demo
ENV note1="The First Note First" note2=The\ Second\ Note\ Second \
note3="The Third Note Third"
ENV changeMe="Old Value"
CMD ["sh"]
```



The ARG instruction

When building Docker images, you may need to use variable data to customize the build. The ARG instruction is the tool to handle that situation. To use it, you add ARG instructions to your Dockerfile, and then when you execute the build command, you pass in the variable data with a --build-arg parameter.

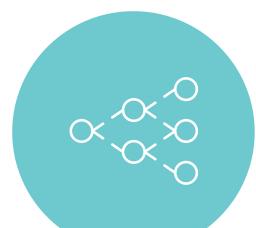
```
→ ~ docker image build --rm \
--build-arg key1="buildTimeValue" \
--build-arg key2="good till env
instruction" \
--tag arg-demo:2.0 .
```

```
# The ARG instruction syntax
ARG <varname>[=<default value>]

# The build-arg parameter syntax
docker image build --build-arg <varname>[=<value>]
...
```

```
# ENV instruction Dockerfile for Docker Quick Start
FROM alpine
LABEL maintainer="Atin <atin@pragra.co>"
ENV key1="ENV is stronger than an ARG"
RUN echo ${key1}
ARG key1="not going to matter"
RUN echo ${key1}

RUN echo ${key2}
ARG key2="defaultValue"
RUN echo ${key2}
ENV key2="ENV value takes over"
RUN echo ${key2}
CMD ["sh"]
```

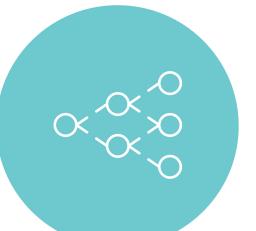


The USER instruction

The USER instruction allows you to set the current user (and group) for all of the instructions that follow in the Dockerfile, and for the containers that are run from the built image. The syntax for the USER instruction is as follows

```
# User instruction syntax  
USER <user>[:<group>] or  
USER <UID>[:<GID>]
```

```
# ENV instruction Dockerfile for Docker  
Quick Start  
FROM alpine  
LABEL maintainer="Atin <atin@pragra.co>"  
RUN id  
USER games:games  
run id  
CMD ["sh"]
```



The WORKDIR instruction

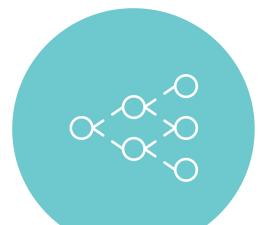
The WORKDIR instruction will change the current working directory in the image to the value provided in the instruction. If any segment of the path in the parameter of the WORKDIR instruction does not yet exist, it will be created as part of the execution of the instruction.

WORKDIR instruction syntax

WORKDIR instruction syntax

WORKDIR /path/to/workdir

```
# WORKDIR instruction Dockerfile for Docker Quick Start
FROM alpine
# Absolute path...
WORKDIR /
# relative path, relative to previous WORKDIR instruction
# creates new folder
WORKDIR sub-folder-level-1
RUN touch file1.txt
# relative path, relative to previous WORKDIR instruction
# creates new folder
WORKDIR sub-folder-level-2
RUN touch file2.txt
# relative path, relative to previous WORKDIR instruction
# creates new folder
WORKDIR sub-folder-level-3
RUN touch file3.txt
# Absolute path, creates three sub folders...
WORKDIR /l1/l2/l3
CMD ["sh"]
```



The VOLUME instruction

The VOLUME instruction will create a storage location that is outside of the United File System, and by so doing, allow storage to persist beyond the life of your container

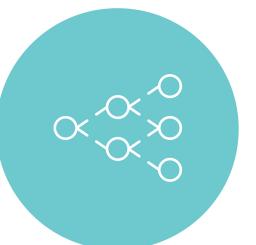
VOLUME instruction syntax

VOLUME ["/data"]

or for creating multiple volumes with a single instruction

VOLUME /var/log /var/db /moreData

```
# VOLUME instruction Dockerfile for Docker
Quick Start
FROM alpine
RUN mkdir /myvol
RUN echo "hello world" > /myvol/greeting
VOLUME /myvol
CMD ["sh"]
```



The EXPOSE instruction

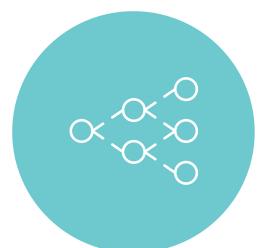
The EXPOSE instruction is a way to document what network ports the image expects to be opened when a container is run from the image built using the Dockerfile.

It is important to understand that including the EXPOSE instruction in the Dockerfile does not actually open network ports in containers. When containers are run from the images with the EXPOSE instruction in their Dockerfile, it is still necessary to include the -p or -P parameters to actually open the network ports to the container.

EXPOSE instruction syntax

EXPOSE <port> [<port>/<protocol>...]

```
# VOLUME instruction Dockerfile for Docker
Quick Start
FROM alpine
RUN mkdir /myvol
RUN echo "hello world" > /myvol/greeting
VOLUME /myvol
CMD ["sh"]
```



The RUN instruction

The RUN instruction is the real workhorse of the Dockerfile. It is the tool by which you affect the most change in the resulting docker image. Basically, it allows you to execute any command in the image. There are two forms of the RUN instruction Every RUN instruction creates a new layer in the image, and the layers for each instruction that follow will be built on the results of the RUN instruction's layer

Exec form of RUN instruction using bash

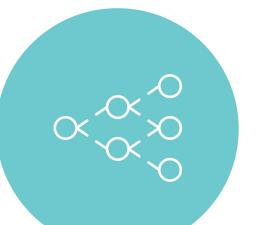
```
RUN ["/bin/bash", "-c", "echo hello world > /myvol/greeting"]
```

RUN <command>

Exec form

```
RUN ["executable", "param1", "param2"]
```

```
# RUN instruction Dockerfile for Docker
Quick Start
FROM ubuntu
RUN useradd --create-home -m -s /bin/bash
dev
RUN mkdir /myvol
RUN echo "hello DQS Guide" > /myvol/greeting
RUN ["chmod", "664", "/myvol/greeting"]
RUN ["chown", "dev:dev", "/myvol/greeting"]
VOLUME /myvol
USER dev
CMD ["/bin/bash"]
```



The CMD instruction

The CMD instruction is used to define the default action taken when containers are run from images built with their Dockerfile. While it is possible to include more than one CMD instruction in a Dockerfile, only the last one will be significant. Essentially, the final CMD instruction provides the default action for the image. This allows you to either override or use the CMD in the image used in the FROM instruction of your Dockerfile.

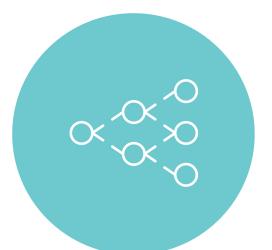
CMD instruction syntax

CMD command param1 param2 (shell form)

CMD ["executable","param1","param2"] (exec form)

**CMD ["param1","param2"] (as default parameters to
ENTRYPOINT)**

```
# RUN instruction Dockerfile for Docker
Quick Start
FROM ubuntu
RUN useradd --create-home -m -s /bin/bash
dev
RUN mkdir /myvol
RUN echo "hello DQS Guide" > /myvol/greeting
RUN ["chmod", "664", "/myvol/greeting"]
RUN ["chown", "dev:dev", "/myvol/greeting"]
VOLUME /myvol
USER dev
CMD ["/bin/bash"]
```



The ENTRYPOINT instruction

The ENTRYPOINT instruction is used to configure a docker image to run like an application or a command. For example, we can use the ENTRYPOINT instruction to make an image that displays help for the curl command.

```
FROM alpine
RUN apk add curl
ENTRYPOINT ["curl"]
CMD ["--help"]
```

We can run the container image with no overriding CMD parameter and it will show help for the curl command.

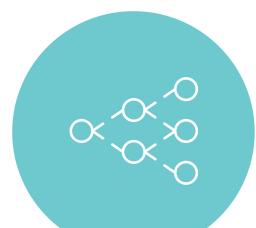
However, when we run the container with a CMD override parameter, in this case, a URL, the response will be to curl the URL

CMD instruction syntax

CMD command param1 param2 (shell form)

CMD ["executable","param1","param2"] (exec form)

**CMD ["param1","param2"] (as default parameters to
ENTRYPOINT)**



The HEALTHCHECK instruction

The HEALTHCHECK instruction, which is a fairly new addition to the Dockerfile, is used to define the command to run inside a container to test the container's application health. When a container has a HEALTHCHECK, it gets a special status variable. Initially, that variable will be set to starting. Any time a HEALTHCHECK is performed successfully, the status will be set to healthy. When a HEALTHCHECK is performed and fails, the fail count value will be incremented and then checked against a retries value. If the fail count equals or exceeds the retries value, the status is set to unhealthy.

```
# HEALTHCHECK instruction Dockerfile for Docker Quick Start
FROM alpine
RUN apk add curl
EXPOSE 80/tcp
HEALTHCHECK --interval=30s --timeout=3s \
CMD curl -f http://localhost/ || exit 1
CMD while true; do echo 'DQS Expose Demo' | nc -l -p 80; done
```

HEALTHCHECK instruction syntax

HEALTHCHECK [OPTIONS] CMD command

HEALTHCHECK NONE (disable any HEALTHCHECK inherited from the base image)

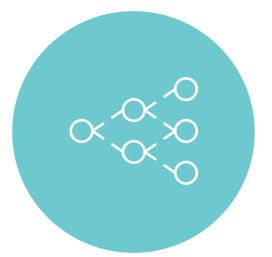
HEALTHCHECK CMD options

--interval=DURATION (default: 30s)

--timeout=DURATION (default: 30s)

--start-period=DURATION (default: 0s)

--retries=N (default: 3)



The ONBUILD instruction

The ONBUILD instruction is a tool used when creating images that will become the parameter to the FROM instructions in another Dockerfile. The ONBUILD instruction just adds metadata to your image, specifically a trigger that is stored in the image and not otherwise used. However, that metadata trigger does get used when your image is supplied as the parameter in the FROM command of another Dockerfile

ONBUILD instruction syntax

ONBUILD [INSTRUCTION]

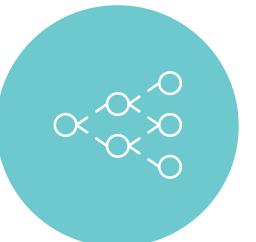
my-base Dockerfile

FROM alpine

LABEL maintainer="Atin <atin@pragra.co>"

ONBUILD LABEL version="1.0"

CMD ["sh"]

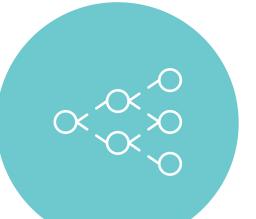


The STOPSIGAL instruction

The STOPSIGAL instruction is used to set the system call signal that will be sent to the container to tell it to exit. The parameter used in the instruction can be an unsigned number, which equals a position in the kernel's syscall table, or it can be an actual signal name in uppercase. Here is the syntax for the instruction:

```
# STOPSIGAL instruction syntax  
STOPSIGAL signal
```

```
# my-base Dockerfile  
FROM alpine  
LABEL maintainer="Atin <atin@pragra.co>"  
ONBUILD LABEL version="1.0"  
CMD ["sh"]
```



Docker image build command

```
# Docker image build command syntax
```

```
Usage: docker image build [OPTIONS] PATH | URL | -
```

```
# Common options used with the image build command
```

```
--rm Remove intermediate containers after a successful build
```

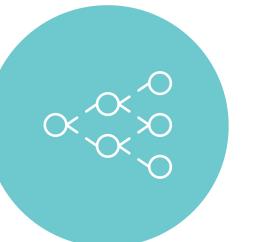
```
--build-arg Set build-time variables
```

```
--tag Name and optionally a tag in the 'name:tag' format
```

```
--file Name of the Dockerfile (Default is 'PATH/Dockerfile')
```

```
→ ~ docker image build --rm --build-arg username=35 --tag arg-demo:2.0 .
```

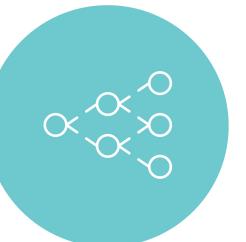
```
→ ~ docker image build --rm -t arg-demo:2.0 .
```



.dockerignore file

The .dockerignore file is used to exclude files that you do not want to be included with the build context during a docker image build. Using it helps to prevent sensitive and other unwanted files from being included in the build context, and potentially in the docker image. It is an excellent tool to help keep your Docker images small.

```
# Example of a .dockerignore file
# Exclude unwanted files
**/*~
**/*.*log
**/.DS_Store
```



Docker image command

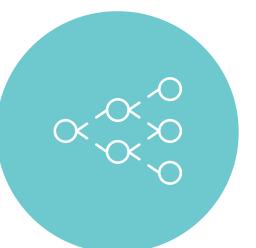
57

Usage: docker image COMMAND

Manage images

Commands:

build	Build an image from a Dockerfile
history	Show the history of an image
import	Import the contents from a tarball to create a
filesystem	image
inspect	Display detailed information on one or more images
load	Load an image from a tar archive or STDIN
ls	List images
prune	Remove unused images
pull	Pull an image or a repository from a registry
push	Push an image or a repository to a registry
rm	Remove one or more images
save	Save one or more images to a tar archive (streamed
to STDOUT by default)	
tag	Create a tag TARGET_IMAGE that refers to
SOURCE_IMAGE	

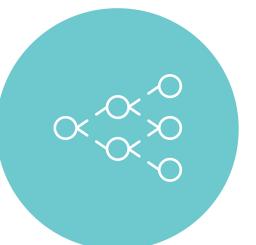


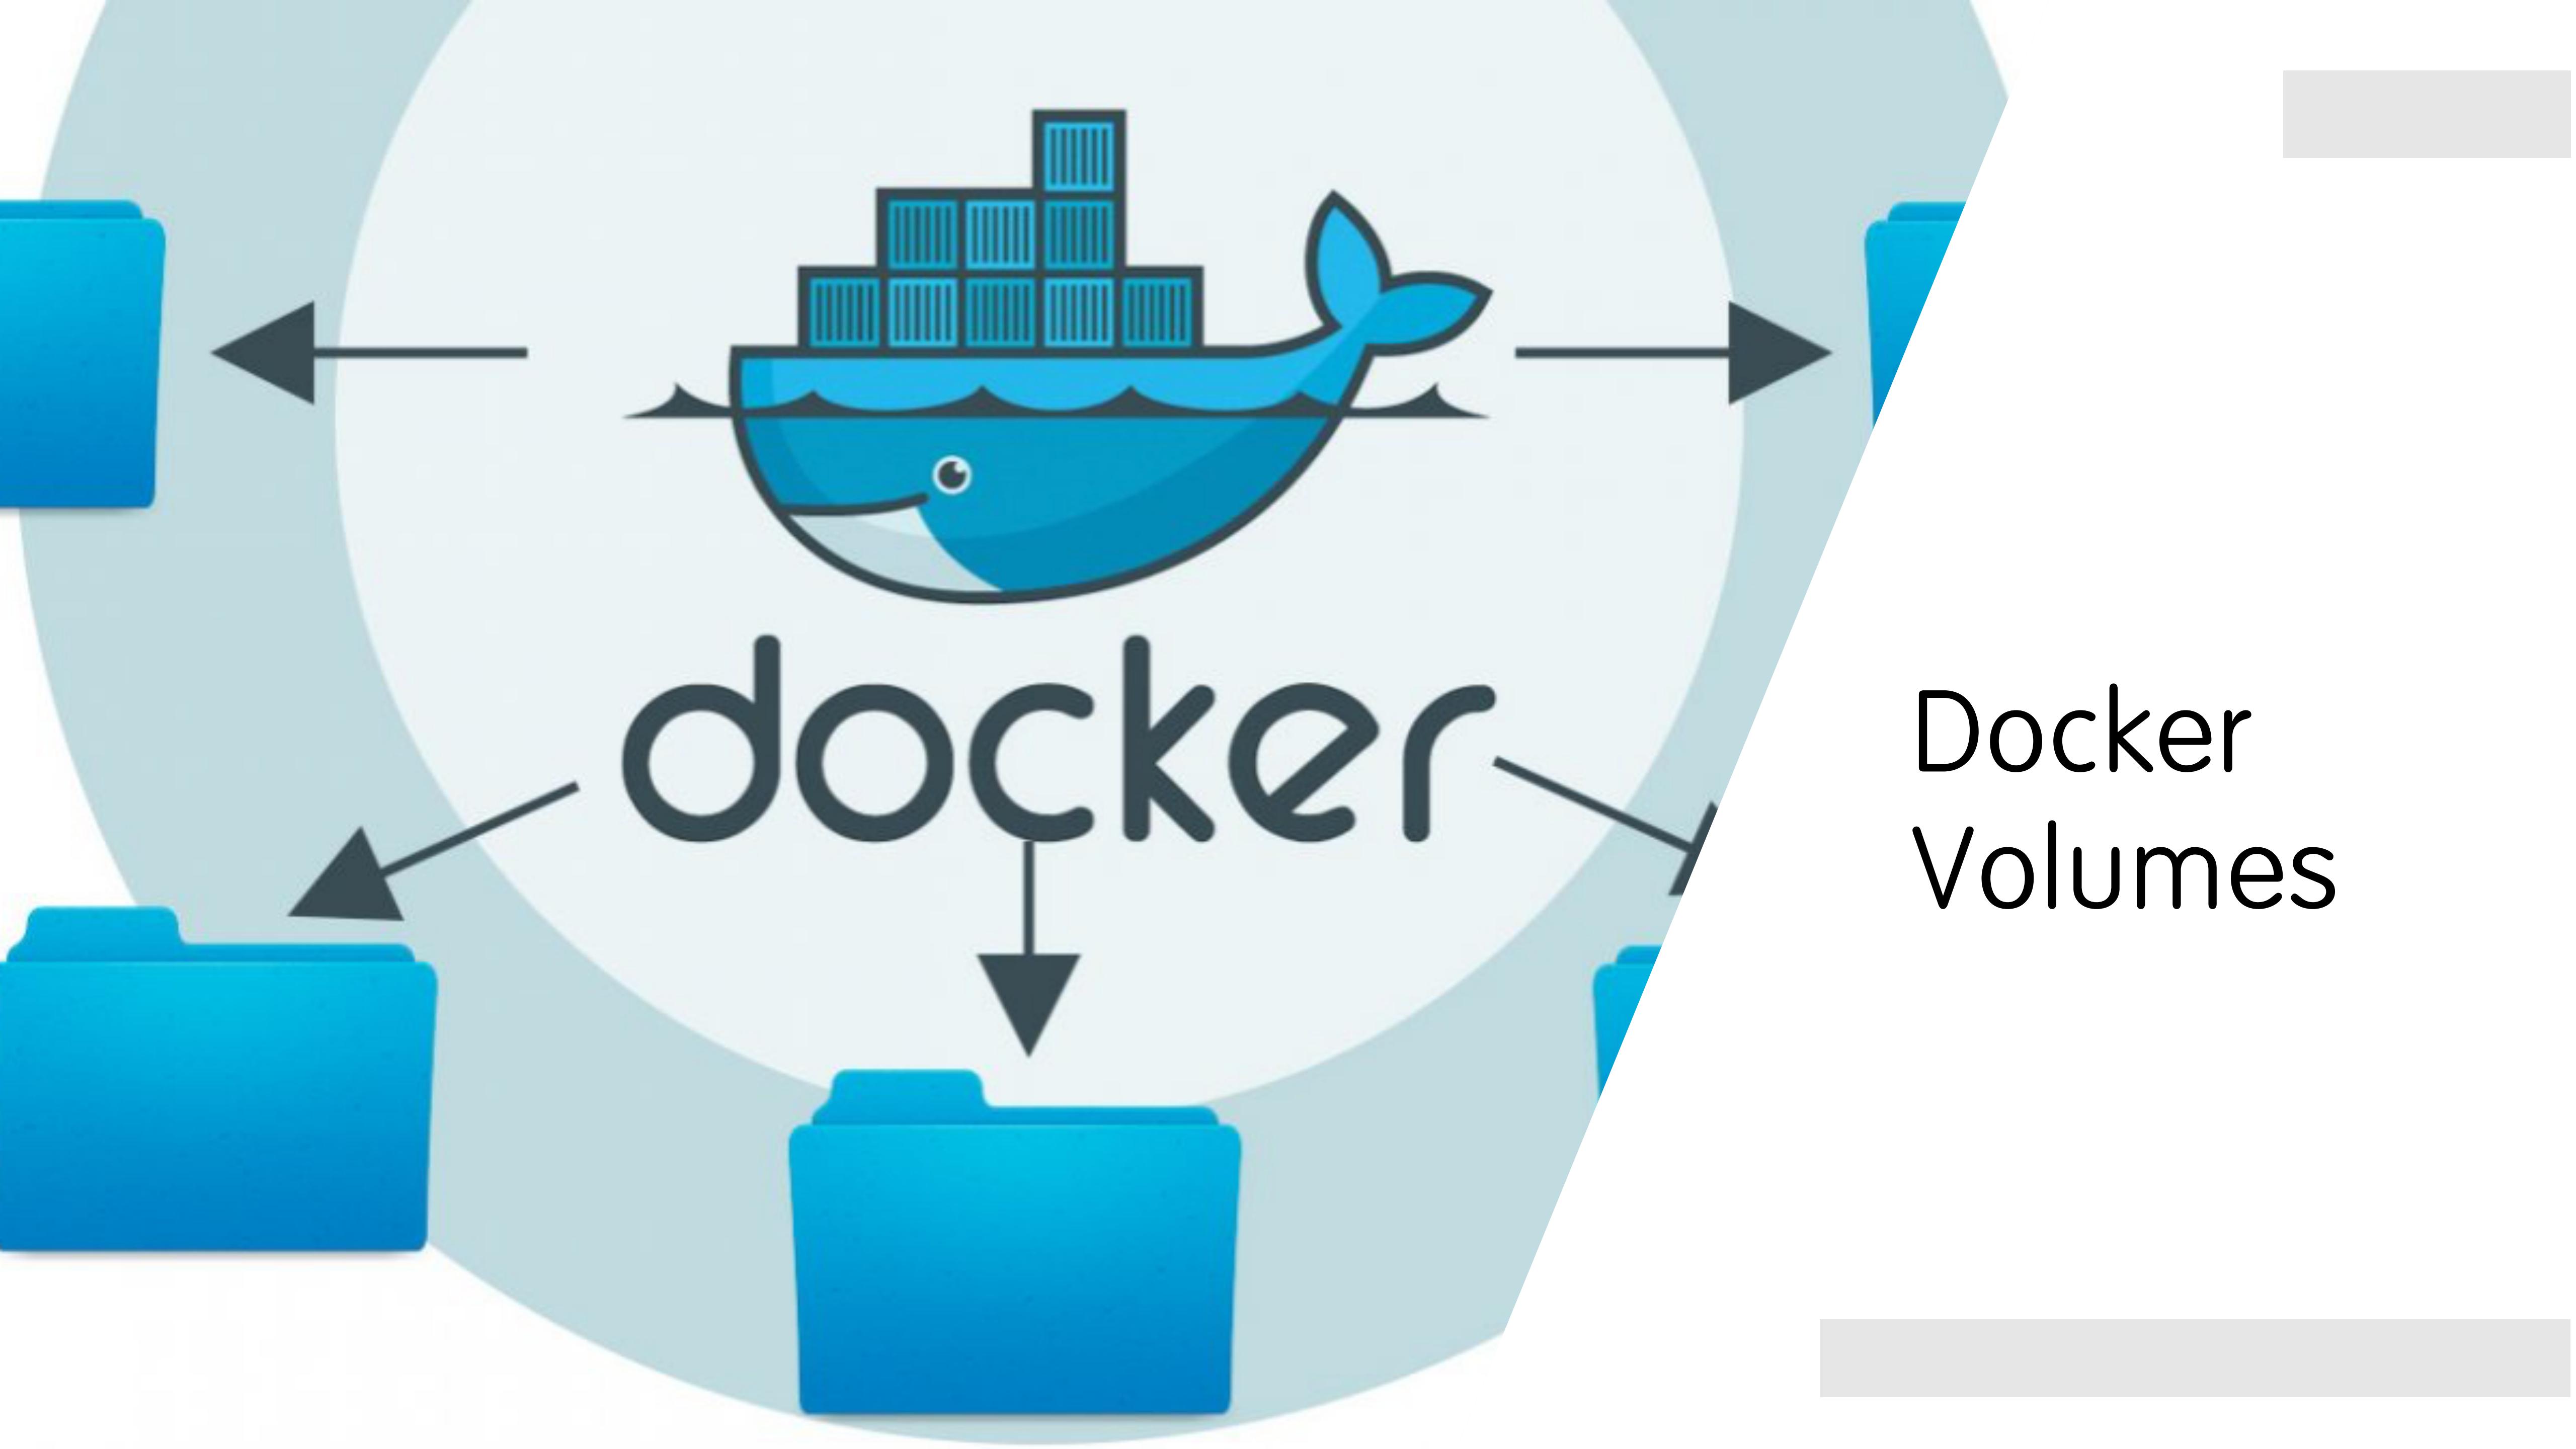
Pragra
ONWARD & UPWARD

Exercise

58

- *Build a docker image from git repo <https://github.com/atinsingh/mongoapi>*
- *Build docker image from alpine linux which has curl*
- *Remove all docker image from your localenv*
- *Create new version of image and tag it with version2.0*
- *Run multiple container from same image and make change to one of the container and save change in image layer.*



The background features the Docker logo, which consists of a stylized blue whale carrying several shipping containers on its back. The word "docker" is written in a lowercase, sans-serif font below the whale. A large, light-blue curved arrow surrounds the logo, pointing clockwise. Four dark-gray arrows point from the corners of the slide towards the center, where the word "docker" is located.

Docker
Volumes

Docker Volumes

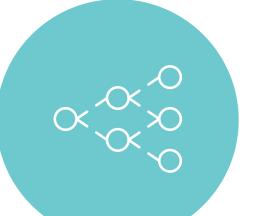
```
# Docker volume management command  
docker volume
```

60

When you run a container from a Docker image, the Docker daemon creates a new read-write layer that holds all of the live data that represents your container. When your container makes changes to its filesystem, those changes go into that read-write layer. As such, when your container goes away, taking the read-write layer goes with it

The Docker volume is a storage location that is completely outside of the Union File System. As such, it is not bound by the same rules that are placed on the read-only layers of an image or the read-write layer of a container. A Docker volume is a storage location that, by default, is on the host that is running the container that uses the volume.

```
# Docker volume management subcommands  
docker volume create # Create a volume  
docker volume inspect # Display information on one or more volumes  
docker volume ls # List volumes  
docker volume rm # Remove one or more volumes  
docker volume prune # Remove all unused local volumes
```



CREATE Volumes

```
# Syntax for the volume create command  
Usage: docker volume create [OPTIONS] [VOLUME]
```

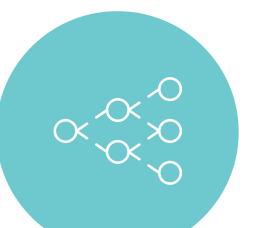
61

When you create a volume default driver is local

```
# The options available to the volume create command:  
-d, --driver string # Specify volume driver name (default "local")  
--label list # Set metadata for a volume  
-o, --opt map # Set driver specific options (default map[])
```

To create volume without specifying the driver and name

```
→ ~ docker volume create  
54f20f4a18ca858fc91c9b309fc713254a1c4b0c33ade516bb737dd6b0285c57  
  
→ ~ docker volume ls  
DRIVER          VOLUME NAME  
local           54f20f4a18ca858fc91c9b309fc713254a1c4b0c33ade516  
bb737dd6b0285c57
```



CREATE Volumes

62

```
# Syntax for the volume create command  
Usage: docker volume create [OPTIONS] [VOLUME]
```

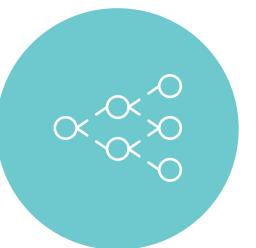
To create volume by specifying name

```
→ ~ docker volume create pragra-volume  
pragra-volume
```

```
→ ~ docker volume ls  
DRIVER          VOLUME NAME  
local           54f20f4a18ca858fc91c9b36bb737dd6b0285c57  
local           pragra-volume
```

Lets inspect a volume

```
→ ~ docker volume inspect pragra-volume  
[  
 {  
   "CreatedAt": "2019-07-24T22:07:17Z",  
   "Driver": "local",  
   "Labels": {},  
   "Mountpoint": "/var/lib/docker/volumes/pragra-volume/_data",  
   "Name": "pragra-volume",  
   "Options": {},  
   "Scope": "local"  
 }  
 ]
```



MOUNT Volume

63

```
# Syntax for the volume create command  
Usage: docker volume create [OPTIONS] [VOLUME]
```

Lets bake a image and mount the volume.

```
→ ~ docker container run --rm -d \  
--mount source=pragra-volume,target=/myvol \  
--name vol-demo2 \  
pragra-vol-demo:1.0 tail -f /dev/null
```

Dockerfile

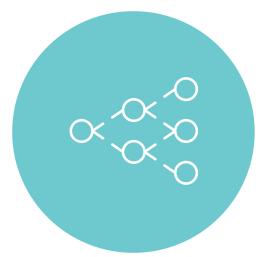
```
FROM alpine  
RUN mkdir /myvol  
RUN echo "Data from image" > /myvol/both-places.txt  
CMD ["sh"]
```

Alternative # Using --mount with source and destination

```
→ ~ docker container run --rm -d \  
--mount source=my-  
volume,destination=/myvol,readonly \  
--name vol-demo2 \  
volume-demo:latest tail -f /dev/null
```

#Using -v

```
→ ~ docker container run --rm -d \  
-v my-volume:/myvol:ro \  
--name vol-demo3 \  
volume-demo:latest tail -f /dev/null
```



Remove Volume

Remove specified volume from system

→ ~ docker volume rm pragra-volume

Remove all volume from system

→ ~ docker volume prune

Find all dangling unused volume

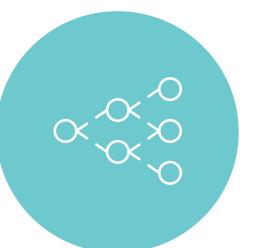
→ ~ docker volume ls --filter
dangling=true

Removed all unused and dangling

~ docker volume rm \$(docker volume ls --filter dangling=true -q)

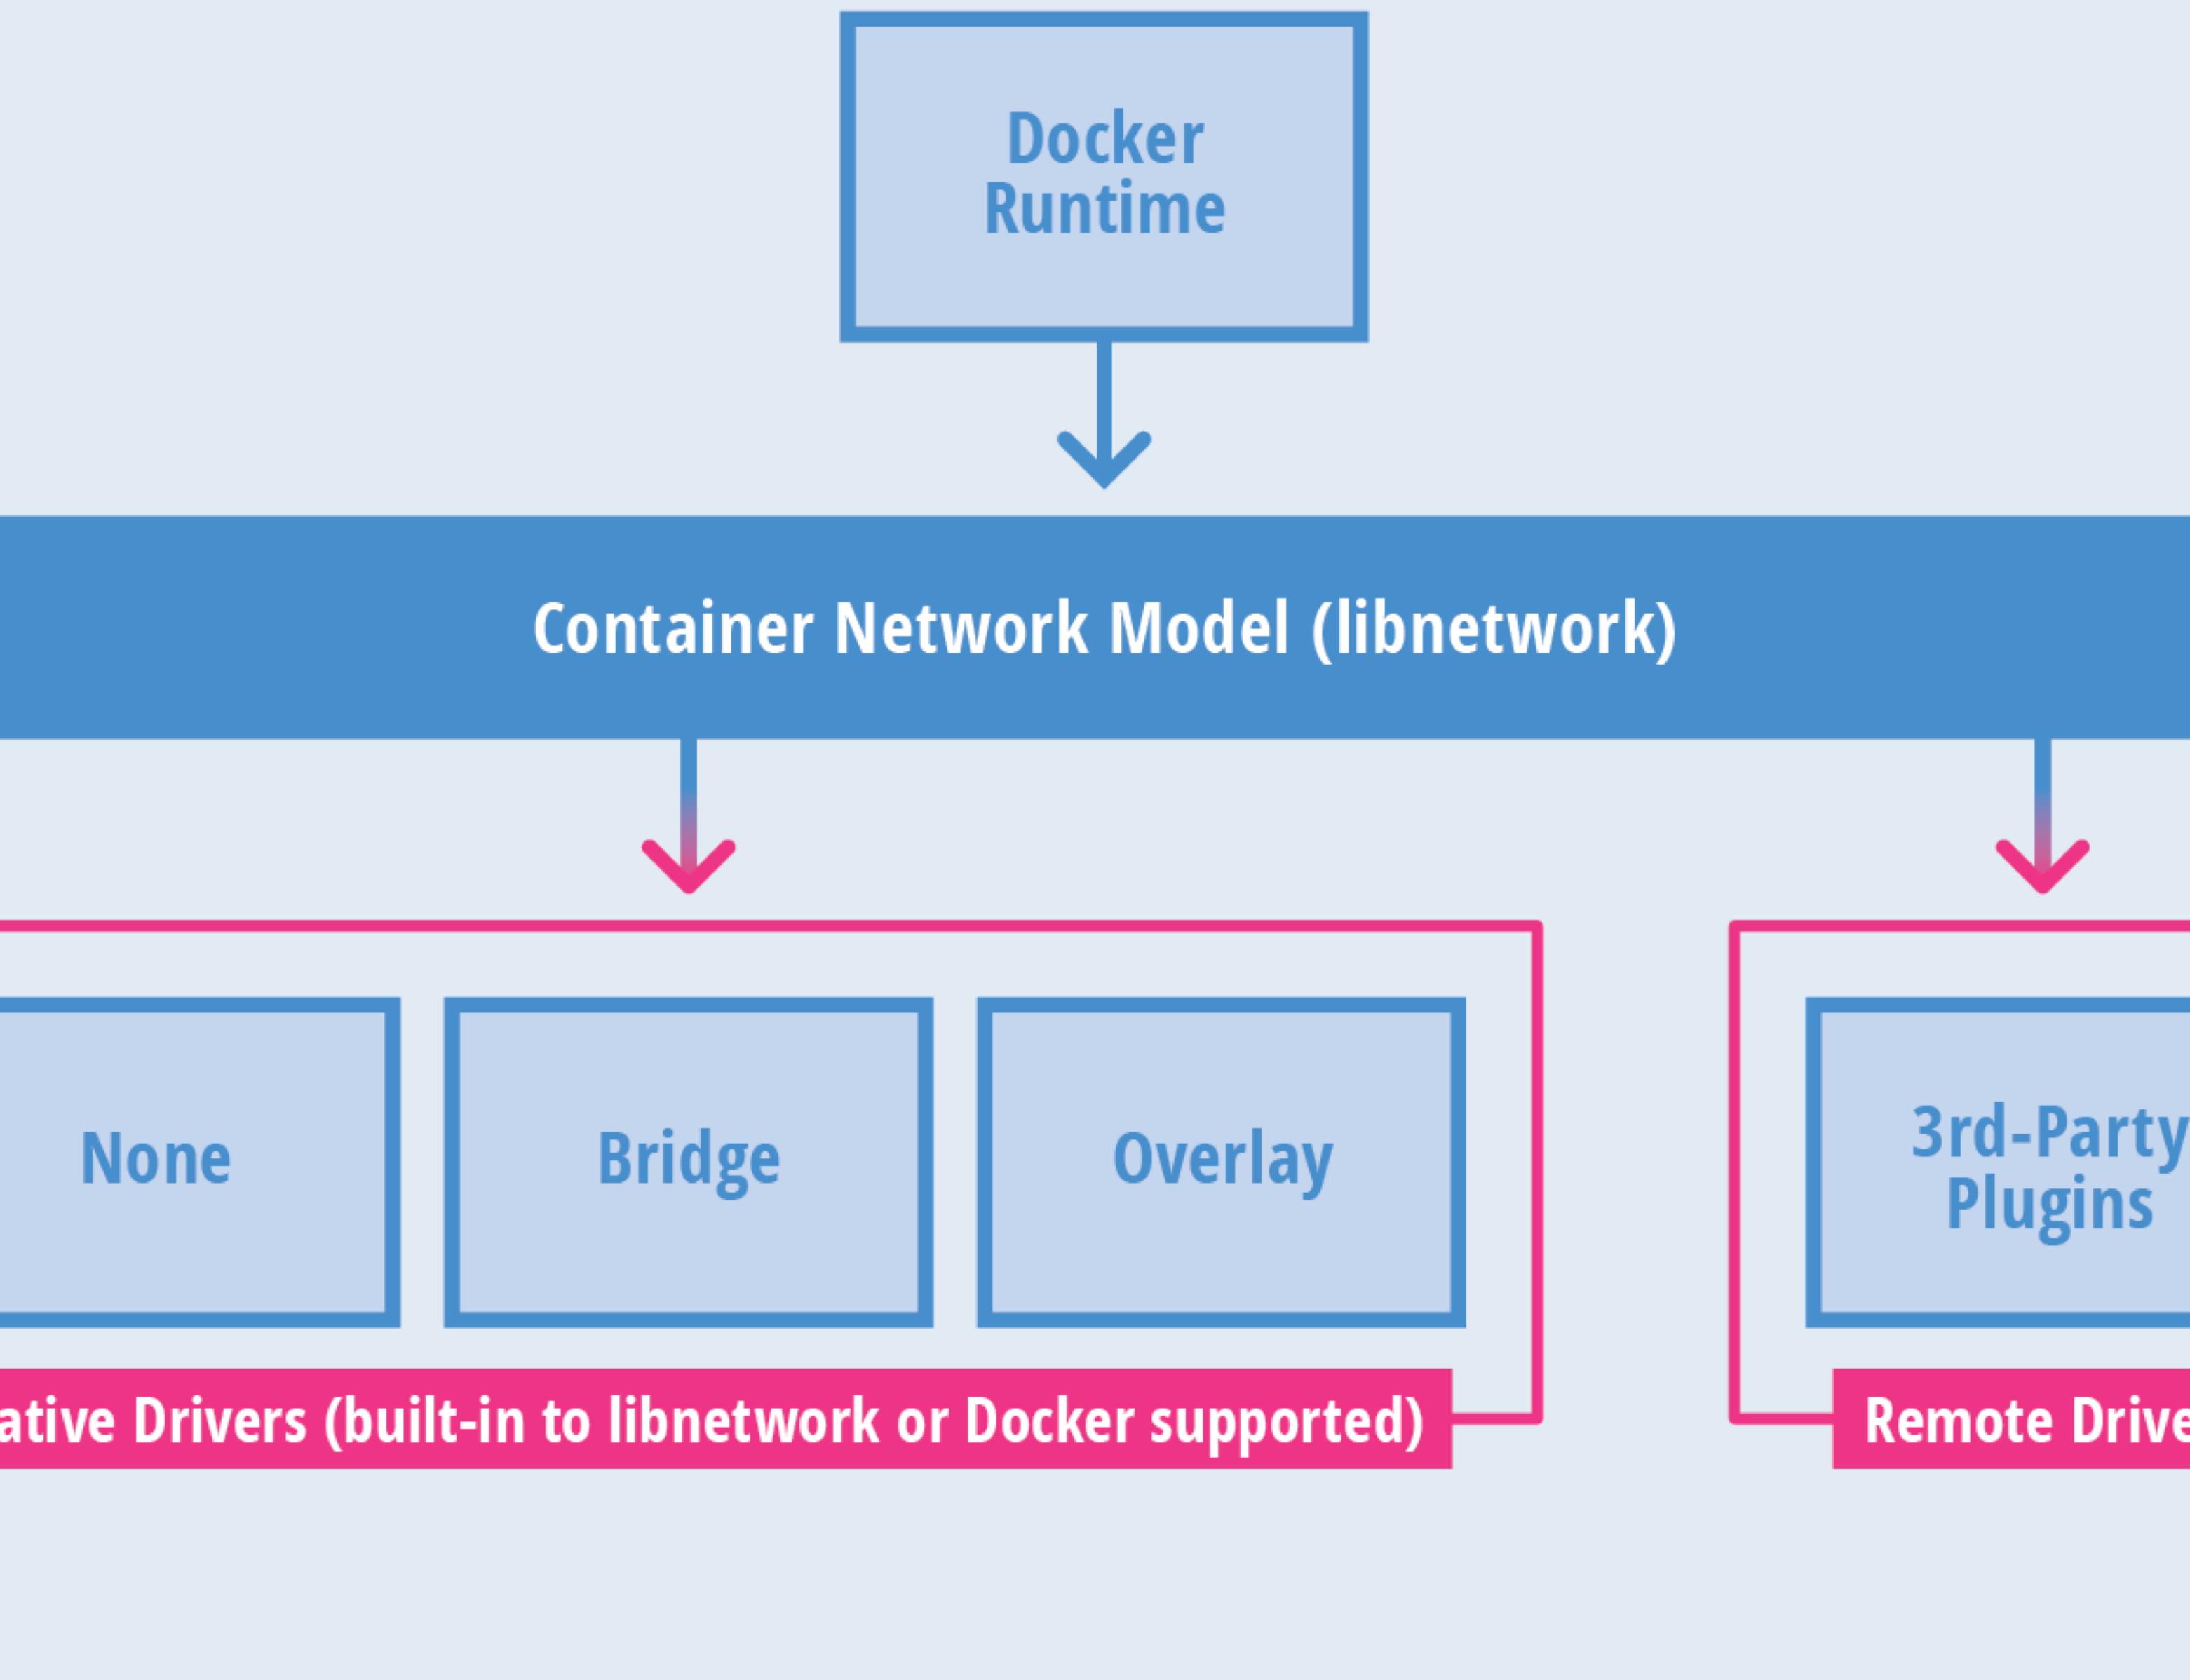
```
# Remove volumes command syntax
Usage: docker volume rm [OPTIONS] VOLUME [VOLUME...]
# Prune volumes command syntax
Usage: docker volume prune [OPTIONS]
```

```
FROM alpine
RUN mkdir /myvol
RUN echo "Data from image" > /myvol/both-places.txt
CMD ["sh"]
```

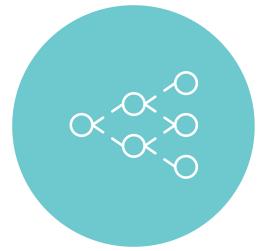


Container Network Model (CNM) Drivers

65



Docker NETWORK



Pragra
ONWARD & UPWARD

◀ *Q&A*

